



UniTs - University of Trieste

Faculty of Data Science and Artificial Intelligence
Department of mathematics informatics and geosciences

Advanced Programming

Lecturer:
Prof. Pasquale Claudio Africa

Authors:

**Christian Faccio
Andrea Spinelli**

February 4, 2025

This document is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike \(CC BY-NC-SA\)](#) license. You may share and adapt this material, provided you give appropriate credit, do not use it for commercial purposes, and distribute your contributions under the same license.

Abstract

As a student of the “Data Science and Artificial Intelligence” master’s degree at the University of Trieste, I have created these notes to study the course “Advanced Programming” held by Prof. Pasquale Claudio Africa. The course aims to provide students with a solid foundation in programming, focusing on the C++ programming language. The course covers the following topics:

- Bash scripting
- C++ basics
- Object-oriented programming
- Templates
- Standard Template Library (STL)
- Libraries
- Makefile
- CMake
- C++11/14/17/20 features
- Integration with Python
- Parallel programming (not covered in the lectures but useful for the HPC course)

While these notes were primarily created for my personal study, they may serve as a valuable resource for fellow students and professionals interested in this field.

Contents

1 Unix Shell	1
1.1 What is a Shell?	1
1.1.1 Types of shell	1
1.2 Shell Scripting	2
1.2.1 Variables and Environmental Variables	2
1.2.2 Initialization Files	3
1.2.3 Basic Shell Commands	3
1.2.4 Shell Scripts	4
1.2.5 Functions	5
1.2.6 Additional Shell Commands	5
1.3 Git introduction	8
1.3.1 SSH Authentication	9
2 C++ Basic	10
2.1 The Birth and Evolution of C++	10
2.1.1 Key Features of C++	10
2.1.2 Modern Applications and Impact	10
2.2 The build process	11
2.2.1 Compiled vs. Interpreted Languages	11
2.2.2 The Build Process	11
2.3 Structure of a Basic C++ Program	13
2.3.1 Overview of Program Structure	13
2.3.2 C++ as a Strongly Typed Language	14
2.3.3 <code>NULL</code> , <code>NaN</code> , and Key Differences	16
2.3.4 Initialization and Aliases	17
2.3.5 The <code>auto</code> Keyword and Type conversion	17
2.4 Memory Management	18
2.4.1 Stack Memory	18
2.4.2 Heap Memory	19
2.4.3 Key Differences Between Stack and Heap	19
2.4.4 When to Use Stack vs. Heap	19
2.4.5 Variables and Pointers	20
2.4.6 Lifetime and Scope	20
2.5 Condition Statements	21
2.5.1 If-Else Statements	21
2.5.2 Switch Statements	21
2.5.3 For loop	21
2.5.4 While Loop	21
2.6 Functions and operators	22
2.6.1 Functions	22
2.7 Operators	22

2.8	User-defined types	23
2.9	Declarations and Definitions	24
2.10	Code Organization	24
2.10.1	Best practices	24
2.11	The Build toolchain in practice	25
2.11.1	Preprocessor and Compiler	25
2.11.2	Linker	25
2.11.3	Preprocessor, Compiler, Linker: Simplified Workflow	26
2.11.4	Loader	26
3	Object-Oriented Programming	27
3.1	Introduction	27
3.2	Classes and Objects in C++	28
3.2.1	Classes and Members	28
3.2.2	Constructors and Destructors	30
3.2.3	The <code>inline</code> Directive	35
3.2.4	In-Class and Out-of-Class Definitions	36
3.2.5	Encapsulation and Access Control	37
3.2.6	Class vs. Struct	38
3.2.7	Getter and Setter Methods	39
3.2.8	Friend Classes	39
3.2.9	Operator Overloading	40
4	Standard Template Library	43
4.1	Containers	43
4.1.1	Sequence Containers	43
4.1.2	Associative Containers	44
4.1.3	Container Adaptors	44
4.1.4	Special Containers	45
4.1.5	Iterators	46
4.1.6	<code>size_type</code> and <code>size_t</code>	48
4.2	Algorithms	48
4.2.1	Non-modifying Algorithms	48
4.2.2	Modifying Algorithms	49
4.2.3	Inserters	50
4.2.4	Sorting Algorithms	51
4.2.5	Min and Max	51
4.2.6	Numeric Algorithms	52
4.3	STL evolution	53
4.4	Smart Pointers	54
4.4.1	<code>std::unique_ptr</code>	55
4.4.2	<code>std::shared_ptr</code>	55
4.4.3	<code>std::weak_ptr</code>	56
4.5	Move Semantics	56
4.5.1	How Move Semantics is implemented	58
4.6	Exceptions	59
4.6.1	Run-time assertions	59
4.6.2	Compile-time assertions	60

4.6.3	Exception handling in C++	60
4.7	STL Utilities	61
4.7.1	I/O Streams	61
4.7.2	Random Numbers	62
4.7.3	Shuffling	63
4.7.4	Sampling	63
4.7.5	Time Measuring	63
4.7.6	Filesystem utilities	64
4.7.7	Filesystem Utilities	64
5	CMake	67
5.1	Introduction	67
5.2	CMakeLists.txt	67
5.2.1	Minimum Version	67
5.2.2	Setting a project	67
5.2.3	Making an executable	68
5.2.4	Making a library	68
5.2.5	Targets	68
5.2.6	Variables	68
5.2.7	Properties	69

1

Unix Shell

1.1 What is a Shell?

The shell is the primary interface between users and the computer's core system. When you type commands in a terminal, the shell interprets these instructions and communicates with the operating system to execute them. It serves as a crucial layer that makes complex system operations accessible through simple text commands.

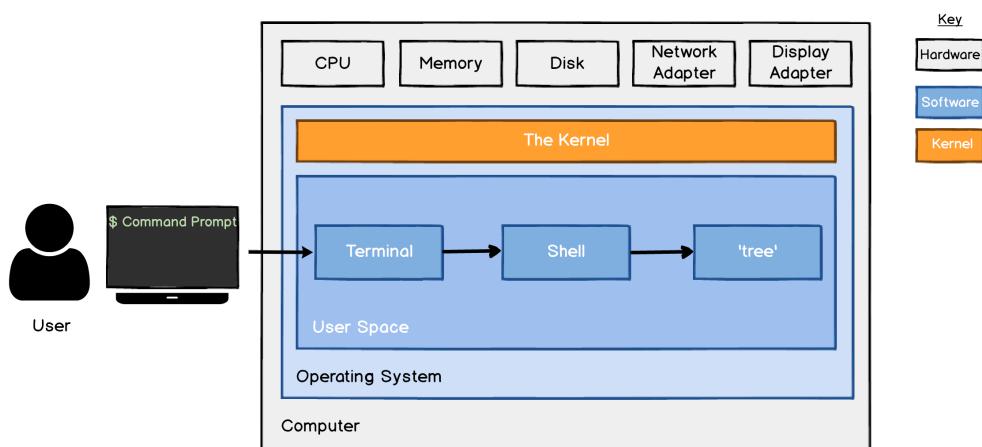


Figure 1.1: The Shell

Definition: Shell

A shell is a program that provides the traditional, text-only user interface for Linux and other UNIX-like operating systems. Its primary function is to read commands that are typed into a console [...] and then execute (i.e., run) them. The term shell derives its name from the fact that it is an outer layer of an operating system. A shell is an interface between the user and the internal parts of the OS (at the very core of which is the kernel).

~Linfo

1.1.1 Types of shell

Over time, various shells have evolved to meet different needs. The most widely used is `Bash` (Bourne Again Shell), named after its creator Stephen Bourne. It's the go-to shell for most Linux systems, serving both as a command interpreter and scripting language. Notably, macOS switched to `zsh` (Bash-compatible) from Catalina onward, while other alternatives include `Fish`, `ksh`, and `PowerShell`.

Tip: Changing the Shell

The shell might be changed by simply typing its name and even the default shell might be changed for all sessions.

Login vs. Non-login Shell

A **login shell** is invoked when a user logs into the system (e.g., through a virtual terminal by pressing `Ctrl+Alt+F1`). It requires the user to provide a username and password. Once authenticated, the user is presented with an interactive shell session. Login shells are typically the first point of interaction between a user and the system.

A **non-login shell**, on the other hand, does not require the user to log in again, as it is executed within an already active user session. For example, opening a graphical terminal in a desktop environment provides a non-login (interactive) shell. Non-login shells are commonly used in environments where the user is already authenticated.

Interactive vs. Non-interactive Shell

An **interactive shell** allows users to type commands and receive immediate feedback. Both login and non-login shells can be interactive. Examples include graphical terminals and virtual terminals. In interactive shells, the prompt (`$PS1`) must be set, which provides the user interface for command input. A **non-interactive shell**, however, is typically executed in automated environments, such as scripts or batch processes. Input and output are generally hidden unless explicitly managed by the calling process. Non-interactive shells are usually non-login shells since the user is already authenticated. For instance, when a script is executed, it runs in a non-interactive shell. However, scripts can emulate interactivity by prompting users for input.

1.2 Shell Scripting

1.2.1 Variables and Environmental Variables

Shells, like any program, use variables to store data. Variables are assigned values using the equals sign (`=`) without spaces. For example, to assign the value `1` to the variable `A`, one would type:

```
A=1
```

To retrieve a variable's value, the dollar sign (`$`) and curly braces are used. For example:

```
echo ${A}
```

Certain variables, called **environmental variables**, influence how processes run on the system. These variables are often predefined. For instance, to display the user's home directory, use: `echo ${HOME}`. To create an environmental variable, prepend the command `export`.

Example: Setting the PATH

For example, to add `/usr/sbin` to the `PATH` environmental variable:

```
export PATH="/usr/sbin:$PATH"
```

The `PATH` variable specifies directories where executable programs are located, ensuring commands can be executed without specifying full paths.

When a terminal is launched, the UNIX system invokes the shell interpreter specified in the `SHELL` environment variable. If `SHELL` is unset, the system default is used. After sourcing initialization files, the shell presents the prompt, which is defined by the `$PS1` environment variable.

1.2.2 Initialization Files

Initialization files are scripts or configuration files executed when a shell session starts. They set up the shell environment, define default settings, and customize behavior. Depending on the type of shell (login, non-login, interactive, or non-interactive), different initialization files are sourced.

- **Login Shell Initialization Files:**

- Bourne-compatible shells: `/etc/profile`, `/etc/profile.d/*`, `~/.profile`.
- Bash: `~/.bash_profile` (or `~/.bash_login`).
- zsh: `/etc/zprofile`, `~/.zprofile`.
- csh: `/etc/csh.login`, `~/.login`.

- **Non-login Shell Initialization Files:**

- Bash: `/etc/bash.bashrc`, `~/.bashrc`.

- **Interactive Shell Initialization Files:**

- `/etc/profile`, `/etc/profile.d/*`, and `~/.profile`.
- For Bash: `/etc/bash.bashrc` and `~/.bashrc`.

- **Non-interactive Shell Initialization Files:**

- For Bash: `/etc/bash.bashrc`.

However, most scripts begin with the condition `[-z "$PS1"] && return`. This means that if the shell is non-interactive (as indicated by the absence of the `$PS1` prompt variable), the script stops execution immediately.

- Depending on the shell, the file specified in the `$ENV` (or `$BASH_ENV`) environment variable may also be read.

1.2.3 Basic Shell Commands

To become familiar with the shell, let's start with some fundamental commands:

- `echo` : Prints the text or variable values you provide at the shell prompt.
- `date` : Displays the current date and time.
- `clear` : Clears the terminal screen.
- `pwd` : Stands for *Print Working Directory*, showing the current directory the shell is operating in. It is also the default location where commands look for files.
- `ls` : Stands for *List*, and lists the contents of the current directory.
- `cd` : Stands for *Change Directory*, and switches the current directory to the specified path.
- `cp` : Stands for *Copy*, and duplicates files or directories from a source to a destination.
- `mv` : Stands for *Move*, and transfers files or directories from one location to another. It can also rename files.
- `touch` : Creates a new, empty file or updates the timestamps of an existing file.
- `mkdir` : Stands for *Make Directory*, and creates new directories.
- `rm` : Stands for *Remove*, and deletes files or directories. To delete directories, the recursive option (`-r`) must be used.

 **Warning: Remove Command**

Be cautious when using the `rm` command, as it permanently deletes files and directories **without moving them to the trash**.

1.2.4 Shell Scripts

Commands can be written in a script file, which is a text file containing instructions for the shell to execute. The first line of the script, known as the **shebang**, specifies the interpreter to use.

```
#!/bin/bash  
#!/usr/bin/env python
```

To make a script executable, you need to change its permissions: `chmod +x script_file`

Not All Commands Are the Same

Commands in the shell can behave differently depending on how they are executed. For example, when a command like `ls` is run, it creates a **subprocess**, a separate instance that inherits the environment of the parent shell. This subprocess runs the command and then terminates, returning control to the parent shell.

💡 Tip: Running Commands

- **Subprocess**: Subprocesses can't modify the parent shell's environment or state. Changes made in subprocesses don't persist.
- **source** : The `source` command (or `.`) runs scripts in the current shell context, allowing environment modifications.
- **Scripts**: Running with `./script_file` creates a subprocess, isolating effects from the parent shell.

Types of Commands

In the shell, commands can fall into several categories:

- **Built-in Commands**: These are commands provided directly by the shell, such as `cd`, that are executed without creating a subprocess, necessary to update environment variables.
- **Executables**: These are standalone programs stored in directories specified by the `$PATH` environment variable. Examples include `ls`, `grep`, and `find`.
- **Functions and Aliases**: These are user-defined commands or shortcuts, often configured in initialization files like `~/.bashrc`.

To determine the type of a command and its location you can use:

- `type command_name` to identify if the command is built-in, an executable, or a function/alias.
- `which command_name` to find the exact location of an executable in the file system.

⚠️ Warning: Spaces in File Names

Avoid using spaces or accented characters in file names. Instead, use:

- `my_file_name` (snake case),
- `myFileName` (camel case),
- `my-file-name` (kebab case).

Spaces complicate scripts and make parsing error-prone.

1.2.5 Functions

A **function** in a shell script is a reusable block of code. The syntax is:

```
1 function_name() {  
2     # Commands to execute  
3 }
```

Scripts or functions can access **input arguments** passed to them using special variables:

- **\$0** : The name of the script or function.
- **\$1, \$2, \$3**, etc.: The first, second, third, etc., arguments.
- **\$#** : The number of arguments passed.
- **\$@** : All arguments as separate words.
- **\$*** : All arguments as a single word (rarely used).

② Example: Function that prints the sum of two numbers

```
1 sum() {  
2     echo $(( $1 + $2 ))  
3 }
```

1.2.6 Additional Shell Commands

More Commands

- **cat** : *Concatenate*. Reads and outputs the contents of files. It can read multiple files and concatenate their content.
- **wc** : *Word Count*. Provides statistics like newline, word and byte count for a list of files.
- **grep** : *Global Regular Expression Print*. Searches for lines containing a specific string or matching a given pattern.
- **head** : Displays the first few lines of a file.
- **tail** : Displays the last few lines of a file.
- **file** : Examines specified files to determine their type.

Redirection, Pipelines, and Filters

Commands can be combined using operators to manipulate input and output streams:

- The **pipe operator** (**|**) forwards the output of one command to another.
Example: `cat /etc/passwd | grep <word>` filters system information for a specific word.
- The **redirect operator** (**>**) sends the standard output to a file.
Example: `ls > files-in-this-folder.txt`.
- The **append operator** (**>>**) appends output to an existing file.
- The **operator** (**&>**) redirects both standard output and standard error to a file.
- **Logical operators**:
 - **&&** : Executes the next command only if the previous one succeeds.
 - **||** : Executes the next command only if the previous one fails.
 - **;** : Executes commands sequentially, regardless of the status of the previous command.
- **\$?** : Contains the exit status of the last command.

Advanced Commands

Cmd	Description	Syntax
tr	<p>Translates characters in <code>SET1</code> to corresponding characters in <code>SET2</code>. If <code>SET2</code> is omitted, deletes characters in <code>SET1</code>.</p> <p>Options:</p> <ul style="list-style-type: none"> • <code>-d</code> : Delete characters in <code>SET1</code>. • <code>-s</code> : Squeeze repeated characters. 	<code>tr [options] SET1 [SET2]</code>
sed	<p>Stream editor for filtering and transforming text.</p> <p>Common uses:</p> <ul style="list-style-type: none"> • Substitution: <code>'s/pattern/replacement/flags'</code> • Deletion: <code>'Nd'</code> deletes the N-th line. <p>Flags:</p> <ul style="list-style-type: none"> • <code>g</code> : Global replacement. • <code>n</code> : Replace n-th occurrence. 	<code>sed [options] 'command' file</code>
cut	<p>Removes sections from each line.</p> <p>Extraction methods:</p> <ul style="list-style-type: none"> • <code>-b list</code>: Select bytes. • <code>-c list</code>: Select characters. • <code>-f list</code>: Select fields; delimiter specified with <code>-d</code>. 	<code>cut [options] file</code>
find	<p>Searches for files in a directory hierarchy.</p> <p>Expressions:</p> <ul style="list-style-type: none"> • <code>-name pattern</code> : Matches filenames. • <code>-type [f d]</code> : Finds files (<code>f</code>) or directories (<code>d</code>). • <code>-exec command {} \;</code> : Executes command on each match. 	<code>find [path ...] [expression]</code>
locate	<p>Searches a database for filenames matching a pattern. Faster than <code>find</code>, but requires database updates.</p> <p>Options:</p> <ul style="list-style-type: none"> • <code>-i</code> : Ignore case. • <code>-r</code> : Use regular expressions. 	<code>locate [options] pattern</code>

?

Example: Usage cases

- **tr :**
 - `echo "hello" | tr 'a-z' 'A-Z'` : Converts lowercase letters to uppercase.
 - `echo "hello123" | tr -d '0-9'` : Removes digits from the input.
- **sed :**
 - `echo 'UNIX is great' | sed 's/UNIX/Linux/'` : Replaces 'UNIX' with 'Linux'.
 - `echo -e "1n2n3" | sed '2d'` : Deletes the second line.
- **cut :**
 - `cut -d ',' -f 1 file.csv` : Extracts the first field from a CSV file.
 - `cut -c 1-5 file.txt` : Extracts the first five characters of each line.

- **find** :
 - `find -type f -name "*.txt"` : Finds all text files in the user's home directory.
 - `find . -type d -name "lib*"` : Finds all directories in the current directory whose names start with "lib".
- **locate** :
 - `locate -i foo` : Finds all files or directories whose name contains 'foo', ignoring case.
 - `locate -r '.*\.txt'` : Finds all files ending with '.txt' using a regular expression.

Quoting in Shell

Quoting affects how strings and variables are interpreted:

- `""` : Double quotes interpret variables.
- `' '` : Single quotes treat everything as literal.

```
1 a=yes
2 echo "$a" # Outputs "yes".
3 echo '$a' # Outputs "$a".
```

The output of a command can be converted into a string and assigned to a variable for later reuse:

```
list=$(ls -l)
# or equivalently:
list='ls -l'
```

Processes

Managing background and foreground processes:

- `./my_command &` : Run a command in the background.
- `Ctrl-Z` : Suspend the current process.
- `jobs` : List background processes.
- `bg %n` : Resume a suspended process in the background.
- `fg %n` : Bring a background process to the foreground.
- `Ctrl-C` : Terminate the foreground process.
- `kill pid` : Send a termination signal to a process.
- `ps aux | grep process` : Find running processes.

Processes in the background are terminated when the terminal closes unless started with `nohup`.

Tip: How to Get Help

- `command -h` or `command --help` : Display a brief help message.
- `man command` : Access the manual for the command.
- `info command` : Show detailed information (if available).

1.3 Git introduction

Git Cheatsheet

Version control is the practice of tracking and managing the changes in the software code. **Git** is an open-source version control system, which helps teams manage changes to the source code over time. It is **distributed**, that is, every developer has the full history of their code repository locally.




Figure 1.2: Git Cheatsheet

Basic Git commands:

- `git init` : Initialize a new Git repository.
- `git clone <url>` : Clone a repository from a URL.
- `git status` : Show the status of the working directory.
- `git add <file>` : Add a file to the staging area.
- `git commit -m "message"` : Commit changes to the repository.
- `git push origin <branch>` : Push changes to a remote repository.
- `git pull origin <branch>` : Pull changes from a remote repository.
- `git branch` : List all branches in the repository.
- `git checkout <branch>` : Switch to a different branch.
- `git merge <branch>` : Merge changes from a branch into the current branch.
- `git log` : Show the commit history.

To collaborate, just remember to pull the changes made by your colleagues before pushing your own changes. If you are working on the same files, the best practice is to create a new **branch**, such that you can merge the changes only when you are sure that everything works as expected.

When you create a branch, Git basically adds a new pointer to the latest commit. This pointer moves forward as you add new commits. When you switch branches, Git moves the pointer to the

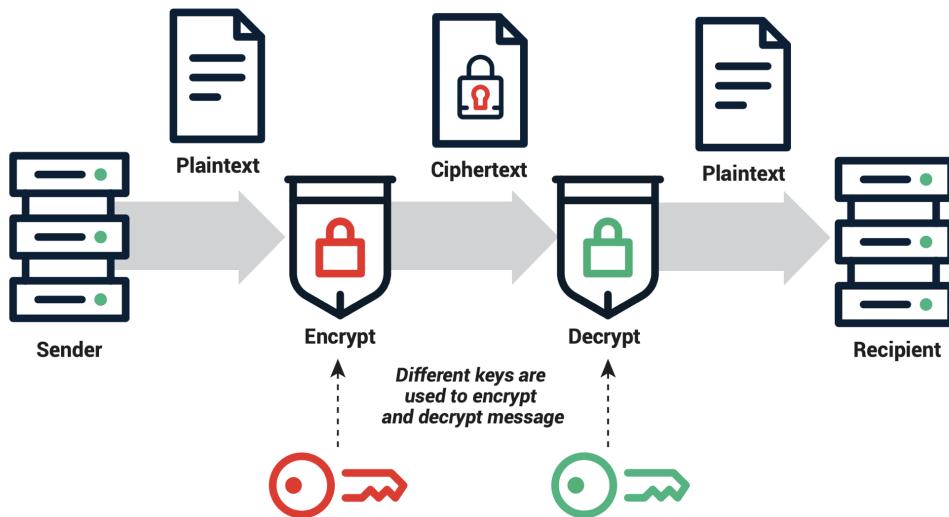


Figure 1.3: SSH

latest commit of the branch you are switching to. When you merge branches, Git combines the changes of the two branches and moves the pointer to the latest commit of the merged branches. If it encounters some conflicts, it will ask you to solve them. Rebasing allow you to move the commits of a branch on top of another branch, so that the history is linear.

Best practices:

- Commit often, with meaningful messages.
- Use branches to isolate changes.
- Pull before pushing.
- Write tests for your code.
- Use a `.gitignore` file to exclude files from version control.

1.3.1 SSH Authentication

To avoid typing your username and password every time you push or pull from a repository, you can use SSH keys for authentication.

Tip: SSH Keys

1. Generate a new SSH key: `ssh-keygen -t rsa -b 4096 -C "youremail"`.
2. Add the SSH key to the SSH agent:
`eval "$(ssh-agent -s)"; ssh-add ~/.ssh/id_rsa`.
3. Add the SSH key to your Git account.
4. Test the SSH connection: `ssh -T`
5. Change the remote URL to the SSH URL: `git remote set-url origin`
6. Test the connection: `git pull`.

2

C++ Basic

2.1 The Birth and Evolution of C++

Programming languages have always evolved to address the growing complexity of software development. Among these, C++ stands out as a language that bridges the gap between low-level system control and high-level abstraction. It combines the performance of C with the principles of object-oriented programming, enabling developers to tackle complex systems efficiently.

The story of C++ begins with C, a powerful yet simple language created by Dennis Ritchie at Bell Labs in the early 1970s. Its efficiency and portability made it a foundation for system programming. In 1979, Bjarne Stroustrup set out to enhance C by introducing support for object-oriented programming (OOP), creating what he called “C with Classes.” This evolved into C++ in 1983, symbolizing an increment over C, and laid the groundwork for modern software engineering. Standardization followed in the late 1980s, ensuring compatibility across platforms. Over the years, successive standards like C++11, C++17, and C++20 have introduced features such as smart pointers, lambda expressions, and modules, keeping C++ at the forefront of programming innovation.

2.1.1 Key Features of C++

- **Object-Oriented Programming (OOP):** C++ introduced core OOP concepts like classes, inheritance, and polymorphism, enabling developers to design modular, reusable, and maintainable software.
- **Generic Programming:** The inclusion of templates brought the power of generic programming, allowing for flexible and reusable data structures and algorithms. Techniques like template metaprogramming extended this capability further.

2.1.2 Modern Applications and Impact

Today, C++ is used across diverse fields, including game development, embedded systems, scientific computing, and finance. Its combination of performance and expressiveness makes it a vital tool for building software that demands both efficiency and scalability.

An active open-source community, including projects like the Boost C++ Libraries, has greatly expanded its capabilities. With ongoing innovations like concepts and modules, C++ continues to adapt to the demands of modern software development, ensuring its relevance for decades to come.

2.2 The build process

2.2.1 Compiled vs. Interpreted Languages

C++ is a **compiled language**, which means that the source code must be translated into machine code before it can be executed. This translation process is performed by a **compiler**, which reads your source code and generates an executable file that can be run on a computer.

In contrast, **interpreted languages** like Python are executed line by line by an **interpreter**. The interpreter reads the code, evaluates it, and executes the corresponding instructions without generating a separate executable file.

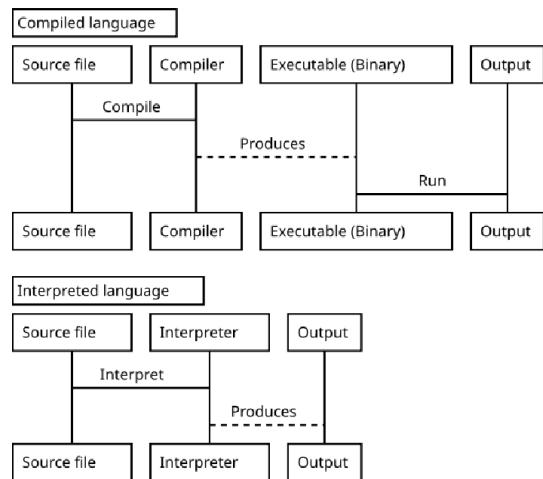


Figure 2.1: Compiled vs. Interpreted Languages

2.2.2 The Build Process

The build process for a C++ program involves several steps:

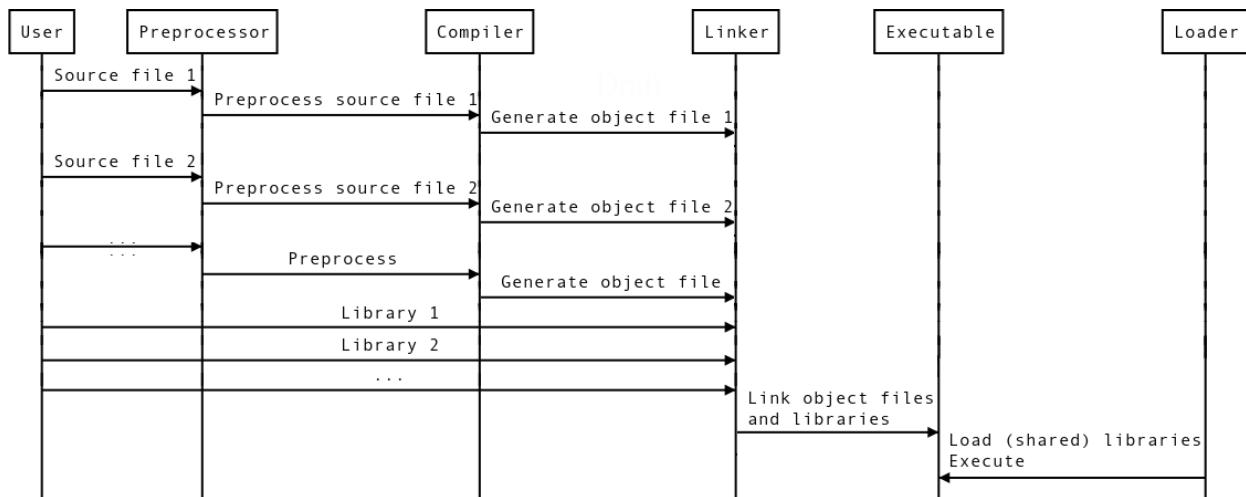


Figure 2.2: The Build Process

Preprocessor

The preprocessor is the first step in the build process. It processes directives that begin with `#` and modifies the source code before it is compiled. Common preprocessor directives include:

- `#include` for including header files;
- `#define` for defining macros;
- `#ifdef`, `#ifndef`, `#else`, `#endif` for conditional compilation;
- `#pragma` for compiler-specific directives.

③ Example: Preprocessor

Original source code:

```
1 #include <iostream>
2 #define GREETING "Hello, World!"
3
4 int main() {
5     std::cout << GREETING << std::endl;
6     return 0;
7 }
```

Preprocessed source code:

```
1 // Content of <iostream>, simplified for demonstration
2 namespace std {
3     extern ostream cout;
4     extern ostream endl;
5 }
6
7 int main() {
8     std::cout << "Hello, World!" << std::endl;
9     return 0;
10 }
```

Compiler

The **compiler** translates the preprocessed source code into assembly or machine code. This phase involves multiple steps:

1. **Lexical Analysis:** Tokenizes the source code into meaningful elements like keywords, identifiers, and operators.
2. **Syntax Analysis (Parsing):** Constructs a syntax tree or abstract syntax tree (AST) to represent the grammatical structure of the code.
3. **Semantic Analysis:** Checks for logical consistency, type compatibility, and adherence to language rules.
4. **Code Generation:** Converts the AST into assembly or machine code.
5. **Optimization:** Enhances the efficiency of the generated code.
6. **Output:** Produces object files containing machine code.

```
g++ main.cpp -o main.o
```

Common compiler options include:

- **-O**: Specify optimization levels (e.g., **-O2**, **-O3**).
- **-g**: Include debugging information for tools like **gdb**.
- **-std**: Specify the C++ standard (e.g., **-std=c++17**, **-std=c++20**).

Linker

The **linker** combines object files into a single executable. This step supports modular programming and ensures that all references between different parts of the program are resolved.

The linking process includes:

1. **Symbol Resolution:** Matches symbols (function, variable names, ...) between object files.
2. **Relocation:** Adjusts memory addresses to create a unified memory layout for the program.
3. **Output:** Produces an executable file.
4. **Linker Errors/Warnings:** Identifies missing symbols or conflicts.

```
1 {Linking Object Files}
2     g++ main.o helper.o -o my_program
```

Linking can be static or dynamic:

- **Static Linking:** All required libraries are included in the final binary, resulting in a larger file size. Libraries do not need to be present on the target system.
- **Dynamic Linking:** Libraries are referenced at runtime, resulting in a smaller binary. Requires the necessary libraries to be present on the system during execution.

Loader

The **loader** prepares the executable for execution by loading it into memory, handling these steps:

1. **Memory Allocation:** Reserves memory for the executable and its data.
2. **Relocation:** Adjusts memory addresses as necessary to account for the executable's location in memory.
3. **Initialization:** Sets up the runtime environment for the program.
4. **Execution:** Begins executing the program's entry point (e.g., `main()` in C++).

⌚ Observation:

Dynamic linking at runtime enhances flexibility by including external libraries only when the program is executed. This approach reduces the initial binary size and allows for library updates without recompiling the application.

2.3 Structure of a Basic C++ Program

2.3.1 Overview of Program Structure

A typical C++ program is composed of a collection of functions. Every C++ program must include the `main()` function, which serves as the entry point. Additional functions can be defined as needed, and their statements are enclosed in curly braces `{}`. Statements are executed sequentially unless control structures like loops or conditionals are applied.

💡 Example: Basic Program Structure

```
1 #include <iostream>
2
3 int main() { // Entry point of the program.
4     std::cout << "Hello, world!" << std::endl;
5     return 0; // Indicates successful execution.
6 }
```

- `#include <iostream>` : Includes the Input/Output stream library.
- `int main()` : The entry point function.

- `std::cout` : Standard output stream for printing to the console.
- `<<` : Stream insertion operator.
- `"Hello, world!"` : The string to print.
- `<< std::endl` : Outputs a newline character and flushes the stream.
- `return 0;` : Indicates successful program termination.

How to Compile and Run

After writing your C++ program, use the GNU C++ compiler (`g++`) to create an executable:

```
g++ hello_world.cpp -o hello_world
```

Execute the compiled program from the terminal, optionally passing command-line arguments:

```
./hello_world [arg1] [arg2] ... [argN]
```

Verify the program's execution status by examining its exit code (0 typically indicates success):

```
echo $?
```

2.3.2 C++ as a Strongly Typed Language

C++ enforces strict type checking during compilation. Variables must be declared with a specific type, ensuring type safety and reducing runtime errors.

It provides various built-in **”Fundamental Types”** to handle data of different kinds and sizes. These types include integers, floating-point numbers, characters, Booleans, and more.

Data Type	Size (Bytes)	<code><cstdint></code>
<code>bool</code>	1	
<code>char</code>	1	
<code>signed char</code>	1	<code>int8_t</code>
<code>unsigned char</code>	1	<code>uint8_t</code>
<code>short</code>	2	<code>int16_t</code>
<code>unsigned short</code>	2	<code>uint16_t</code>
<code>int</code>	4	<code>int32_t</code>
<code>unsigned int</code>	4	<code>uint32_t</code>
<code>long int</code>	4 or 8	<code>int32_t</code> or <code>int64_t</code>
<code>long unsigned int</code>	4 or 8	<code>uint32_t</code> or <code>uint64_t</code>
<code>long long int</code>	8	<code>int64_t</code>
<code>long long unsigned int</code>	8	<code>uint64_t</code>
<code>float</code>	4	
<code>double</code>	8	

Table 2.1: Sizes of Fundamental Types in C++

② Example: Strong Typing in C++

```
1 int x = 5;
2 char ch = 'A';
3 float f = 3.14;
4
5 x = 1.6;      // Legal, but truncated to 1.
6 f = "a string"; // Illegal.
7
8 unsigned int y{3.0}; // Uniform initialization: illegal.
```

C++ supports several **integer types** with varying sizes and value ranges. Common types include `int`, `short`, `long`, and `long long`.

```
int age = 30;
short population = 32000;
long long large_number = 123456789012345;
```

Floating-point types represent real numbers. These include `float`, `double`, and `long double`. They are ideal for representing decimal values.

```
float pi = 3.14;
double gravity = 9.81;
```

Floating-Point Arithmetic

Floating-point numbers in C++ are represented using the format $\pm f \cdot 2^e$, where:

- f : the *significand* or *mantissa*, representing the precision of the number.
- e : the *exponent*, determining the scale of the number.
- 2: the base, as floating-point numbers are typically stored in binary form.

This representation enables efficient handling of very large or very small values but comes with certain limitations, such as rounding errors.

② Example:

```
1 double epsilon = 1.0;
2
3 while (1.0 + epsilon != 1.0) {
4     epsilon /= 2.0; // Finding machine epsilon.
5 }

1 double a = 0.1, b = 0.2, c = 0.3;
2
3 if (a + b == c) { // Unsafe comparison.
4     // Due to precision limitations, this might not hold true.
5 }
6
7 if (std::abs((a + b) - c) < 1e-9) {
8     // Use a tolerance for safe comparison.
9 }
```

Normalized Numbers: In normalized form, the most significant bit of the significand is always 1, which ensures efficient use of available precision and avoids redundant representations.

IEEE 754 Standard: The IEEE 754 Standard defines how floating-point numbers are represented and manipulated. It specifies:

- Standardized formats for `float`, `double`, and `long double`.
- Rounding rules to maintain accuracy.
- Special values such as `NaN` (Not-a-Number) and infinity for handling exceptional cases.

Characters and Strings

Characters in C++ are represented using the `char` type, while strings are sequences of characters represented by the `std::string` class.

```
char letter = 'A'; // Single character.  
std::string name = "Alice"; // String of characters.  
std::string greeting = "Hello, " + name + "!"; // Concatenation.
```

Boolean Types

C++ provides a built-in `bool` type for logical values. A `bool` variable can hold one of two values, `true` or `false`, useful for conditional statements and logical operations.

```
bool flag = true;  
if (flag) {  
    std::cout << "Flag is true" << std::endl;  
}
```

Note that numbers can be implicitly converted to `bool`: 0 is `false`, while any other value is `true`.

```
if (0) // false  
if (42) // true
```

2.3.3 `NULL`, `NaN`, and Key Differences

NULL

`NULL` is used to represent a null pointer or an invalid memory address. Its representation depends on the system and context:

- In C/C++, `NULL` is typically defined as `0` or `(void*)0`.
- In memory, a null pointer is often represented by a sequence of **all-zero bits**. (e.g. 32-bit system: `0x00000000`; 64-bit system: `0x0000000000000000`)

Example: Representation of NULL

```
1 int* ptr = NULL; // ptr points to memory address 0x00000000  
2 if (ptr == NULL) {  
3     printf("Pointer is NULL\n");  
4 }
```

NaN

NaN (Not a Number) is a special value used in floating-point arithmetic to represent undefined or unrepresentable results. Its representation is defined by the **IEEE 754 floating-point standard**:

- A NaN value is represented by an **exponent filled with 1s** (all bits set to 1) and a **non-zero significand**.
- There are two types of NaN, qNaN and sNaN:
 - Quiet NaN: Used for undefined or unrepresentable results. It has a leading 1 in the significand.
 - Signaling NaN: Used to trigger exceptions in certain operations. It has a leading 0 in the significand.

⌚ Example: Representation of NaN

```
1 double x = 0.0 / 0.0; // Creates NaN
2 if (isnan(x)) {
3     std::cout << "x is NaN\n";
4 }
```

Key Differences

Feature	NULL	NaN
Purpose	Represents a null pointer	Represents an undefined floating-point value
Data Type	Used with pointers	Used with floating-point numbers
Representation	All-zero bits	Exponent filled with 1s and non-zero significand
Context	Memory addresses, pointers	Floating-point arithmetic

2.3.4 Initialization and Aliases

Initialization assigns an initial value to a variable at the time of declaration. C++ supports several initialization methods: direct, copy, and uniform initialization.

```
int x = 42;           // Direct initialization.
int y(30);           // Constructor-style initialization.
int z{15};            // Uniform initialization (preferred).
```

Type Aliases can create alternative names for existing types using `using` or `typedef`.

```
using integer = int;    // Alias for int.
typedef float distance; // Alias for float.
```

2.3.5 The `auto` Keyword and Type conversion

The `auto` keyword allows the compiler to deduce the type of a variable based on its initialization value. This is useful for simplifying code and avoiding verbose type declarations.

```
auto a{42};           // int.
auto b{12L};           // long.
auto c{5.0F};          // float.
auto d{10.0};           // double.
auto e{false};          // bool.
auto f{"string"};       // char[7].
```

💡 Tip: Best Practices

- Use `auto` for complex types or when the exact type is unimportant.
- Avoid `auto` for publicly visible variables or ambiguous initializations.

C++ supports **implicit and explicit type conversions** to convert between different data types. Implicit conversions are performed automatically by the compiler, while explicit conversions require manual intervention.

❓ Example: Type Conversion

Implicit conversion:

```
1 int x = 10;
2 double y = x; // int to double (implicit).
```

Explicit conversion:

```
1 double z = 3.14;
2 int w = static_cast<int>(z); // double to int (explicit).
```

2.4 Memory Management

In computer programming, memory is divided into two main regions: the **stack** and the **heap**. These regions serve different purposes and are managed differently. Below is a detailed explanation of their differences.

2.4.1 Stack Memory

The **stack** is used for **static memory allocation**, where memory is allocated and deallocated in a last-in, first-out (LIFO) order.

Key Characteristics of Stack Memory

- **Purpose:** Used for storing local variables, function parameters, and return addresses.
- **Management:** Memory allocation and deallocation are handled automatically by the compiler.
- **Speed:** Accessing the stack is very fast because it uses a simple pointer-based mechanism (the stack pointer).
- **Size:** The stack has a limited size, which is determined at the start of the program. If the stack exceeds its size, a **stack overflow** occurs.
- **Lifetime:** Memory is automatically freed when the function or block that allocated it exits.
- **Fragmentation-Free:** The stack does not suffer from fragmentation because memory is always allocated and freed in a strict order.

❓ Example: Stack Memory Example

```
1 void foo() {
2     int x = 10; // 'x' is allocated on the stack
3     // Memory for 'x' is automatically freed when 'foo' exits
4 }
```

2.4.2 Heap Memory

The **heap** is used for **dynamic allocation**, where memory can be allocated and deallocated in any order.

Key Characteristics of Heap Memory

- **Purpose:** Used for dynamically allocated data (e.g., arrays, objects, or data structures whose size is not known at compile time).
- **Management:** Memory allocation and deallocation are managed manually by the programmer (e.g., using `malloc / free` in C or `new / delete` in C++).
- **Speed:** Accessing the heap is slower than the stack as it involves complex memory management.
- **Size:** The heap is much larger than the stack and can grow dynamically as needed (limited only by the system's available memory).
- **Lifetime:** Memory remains allocated until it is explicitly freed by the programmer. If not freed, it leads to **memory leaks**.
- **Fragmentation:** The heap can suffer from fragmentation over time, as memory is allocated and freed in arbitrary order.

③ Example: *Heap Memory Example*

```
1 void bar() {  
2     int* ptr = new int(20); // 'ptr' points to memory allocated on  
// the heap  
3     // Memory for 'ptr' must be explicitly freed  
4     delete ptr; // Free the memory to avoid a memory leak  
5 }
```

2.4.3 Key Differences Between Stack and Heap

Feature	Stack Memory	Heap Memory
Purpose	Static memory allocation	Dynamic memory allocation
Management	Automatic (compiler-managed)	Manual (programmer-managed)
Speed	Very fast	Slower
Size	Limited (predefined size)	Large (limited by system memory)
Lifetime	Automatically freed when scope ends	Must be explicitly freed
Fragmentation	No fragmentation	Can suffer from fragmentation
Usage	Local variables, function parameters	Dynamically allocated data

2.4.4 When to Use Stack vs. Heap

- **Use the Stack:**
 - For small, short-lived data (e.g., local variables, function parameters).
 - When the size of the data is known at compile time.
 - When you want fast and automatic memory management.
- **Use the Heap:**
 - For large or dynamically sized data (e.g., arrays, objects).
 - When the lifetime of the data extends beyond the current scope.
 - When you need flexibility in memory allocation and deallocation.

2.4.5 Variables and Pointers

Variables represent named memory locations, while pointers store memory addresses, enabling direct access and manipulation. **Stack Variables** are declared locally within functions or blocks, are stored on the stack and are accessible directly. **Heap Variables** require pointers for access and explicit deallocation.

⌚ Example: Working with Variables and Pointers

```
1 int value = 10;           // Stack variable
2 int* ptr = &value;        // Pointer to stack variable
3
4 int* heap_var = new int(25); // Pointer to heap-allocated variable
5 *heap_var = 30;           // Modify heap variable through pointer
6
7 // Cleanup
8 delete heap_var;         // Deallocate heap memory
9 heap_var = nullptr;       // Reset pointer.
```

2.4.6 Lifetime and Scope

The lifetime of a variable refers to the duration it exists in memory, while its scope defines where it is accessible in code.

Stack Variables

- Limited to the scope of their defining function or block.
- Automatically deallocated when the scope ends.

Heap Variables

- Persist beyond their defining scope until explicitly deallocated.
- Risk memory leaks if not deallocated properly.

⌚ Example: Lifetime and Scope Example

```
1 void example() {                      // Lifetime:
2     int stack_var = 5;                 // ends with the function.
3     int* heap_var = new int(10);      // persists until delete.
4
5     delete heap_var;                 // Deallocate heap memory.
6     heap_var = nullptr;              // Reset pointer.
7 }
```

💡 Tip: Best Practices for Memory Management

- Use stack memory for small, short-lived variables.
- Use heap memory for large or long-lived data.
- Always match `new` with `delete` and `new[]` with `delete[]`.
- Prefer modern alternatives like `std::unique_ptr` or `std::shared_ptr` to manage heap memory safely.

2.5 Condition Statements

2.5.1 If-Else Statements

The `if-else` statement is used to execute code based on a condition. If the condition is true, the code inside the `if` block is executed; otherwise, the code inside the `else` block is executed.

```
1 int x = 15;
2 if (x != 10) {
3     std::cout << "x is not equal to 10" << std::endl;
4 } else {
5     std::cout << "x is equal to 10" << std::endl;
6 }
```

2.5.2 Switch Statements

The `switch` statement executes different code blocks based on the value of an expression. When a matching case is found, its code block is executed.

```
1 int choice = 2;
2 switch (choice) {
3     case 1:
4         // Code for first option
5         break;
6     case 2:
7         // Code for second option
8         break;
9     ...
10    default:
11        // Code for when no cases match
12        break;
13 }
```

2.5.3 For loop

The `for` loop is used to execute a block of code a specified number of times. It consists of three parts: initialization, condition, and increment/decrement.

```
1 for (int i = 0; i < 5; i++) {
2     std::cout << i << std::endl;
3 }
```

2.5.4 While Loop

The `while` loop is used to execute a block of code as long as a specified condition is true.

```
1 int i = 0;
2 while (i < 5) {
3     std::cout << i << std::endl;
4     i++;
5 }
```

2.6 Functions and operators

2.6.1 Functions

Functions are blocks of code that perform a specific task. They are defined with a return type, name, parameters, and a body. Functions can be called from other parts of the program to execute the code inside them.

```
int add(int a, int b) {
    return a + b;
}
```

If a function does not output a value, its return type is `void`. A function can also output a general `auto` type, which allows the compiler to deduce the return type based on the return statement. Passing parameters by **value** creates a copy of the argument, while passing by **reference** allows the function to modify the original argument. The function can also return pointers or references, enabling the caller to access or modify the original data.

```
void increment(int& x) {
    x++;
}
auto add(int a, int b) {
    return a + b;
}
```

Warning: *const correctness*

`const` correctness is essential for writing safe and maintainable code. It prevents unintended modifications by marking parameters as read-only when they should not be altered.

- **Avoids unintended changes:** Prevents accidental modifications, improving safety.
- **Improves readability:** Clearly expresses the intent of function parameters.
- **Enables optimizations:** Helps the compiler optimize code by ensuring immutability.

```
void print(const std::string& message) {
    std::cout << message << std::endl;
}
```

2.7 Operators

Increment and Decrement Operators

1. **Increment Operator (`++`):** Increases the value of a variable by 1.
2. **Decrement Operator (`--`):** Decreases the value of a variable by 1.

```
1 int x = 5;
2 x++; // x is now 6
3 x--; // x is now 5
```

Function Overloading

Function overloading is a feature in C++ that allows you to define multiple functions with the same name but different parameters. The compiler selects the appropriate function based on the number or types of arguments during the function call.

```
1 int add(int a, int b) {  
2     return a + b;  
3 }  
4 int add(double a, double b) {  
5     return a + b;  
6 }
```

2.8 User-defined types

- **enum:** Defines a set of named constant values.

```
1 enum Color { RED, GREEN, BLUE };  
2 Color c = RED;
```

- **union:** Allows multiple data members to share the same memory location.

```
1 union Data {  
2     int x;  
3     float y;  
4 };  
5 Data d;  
6 d.x = 10;  
7 d.y = 3.14; // Overwrites 'x'
```

- **struct:** Groups related data members into a single unit.

```
1 struct Point {  
2     int x;  
3     int y;  
4 };  
5 Point p = {10, 20};
```

- **class:** Similar to a struct but with additional features like access control (public, private, protected).

```
1 class Circle {  
2 public:  
3     double radius;  
4     double area() {  
5         return 3.14 * radius * radius;  
6     }  
7 };  
8 Circle c;  
9 c.radius = 5.0;  
10 double a = c.area();
```

2.9 Declarations and Definitions

- **Declaration:** Introduces the name and type of a variable, function, or class without allocating memory or defining its implementation.

```
// Declaration
extern int x;
void foo();
```

- **Definition:** Allocates memory and provides the implementation details for a variable, function, or class.

```
// Definition
int x = 10;
void foo() {
    std::cout << "Hello, world!" << std::endl;
}
```

2.10 Code Organization

Modular programming is a software design technique that divides code into separate modules, each responsible for a specific functionality. This approach enhances code readability, maintainability, and reusability. C++ modules can be organized into header files (`.h`), for declarations, and source files (`.cpp`), for definitions.

② Example: *Header and Source Files*

```
1 // module.h (Header file)
2 #ifndef MODULE_H
3 #define MODULE_H
4
5 void foo();
6
7 #endif
8
9 // module.cpp (Source file)
10 #include "module.h"
11
12 void foo() {
13     std::cout << "Hello, world!" << std::endl;
14 }
```

2.10.1 Best practices

- **Header Guards:** Prevent multiple inclusions of the same header file.
- **Include Guards:** Use `#pragma once` or `#ifndef` guards to avoid redundant includes.
- **Separation of Concerns:** Keep declarations in header files and definitions in source files.
- **Forward Declarations:** Use forward declarations to reduce dependencies and improve compilation times.
- **Namespace Usage:** Enclose code in namespaces to prevent naming conflicts.

③ Example: *namespace usage*

```
1 namespace math {
2     int add(int a, int b) {
3         return a + b;
4     }
5 }
6
7 int main() {
8     int sum = math::add(10, 20);
9     return 0;
10 }
```

2.11 The Build toolchain in practice

2.11.1 Preprocessor and Compiler

The build process starts with the **preprocessor** (`cpp`) and the **compiler** (`g++`, `clang++`):

- **Preprocessor:** Handles directives like `#include`, performs macro substitution, and prepares code.
- **Compiler:** Translates preprocessed code into machine-readable object files.

These steps are often combined when using a compiler command like `g++` or `clang++`.

For a project with three files (`module.hpp`, `module.cpp`, `main.cpp`), the following commands illustrate preprocessing and compilation:

```
# Preprocessor step.
g++ -E module.cpp -I/path/to/include/dir -o module_preprocessed.cpp
g++ -E main.cpp -I/path/to/include/dir -o main_preprocessed.cpp

# Compilation step.
g++ -c module_preprocessed.cpp -o module.o
g++ -c main_preprocessed.cpp -o main.o
```

2.11.2 Linker

The **linker** (`ld`) combines object files into an executable program by resolving external references between them. It also links external libraries if required.

```
g++ module.o main.o -o my_program
```

Linking Against Libraries

To link against external libraries, use the `-l` flag for library names (without the `lib` prefix or file extension) and the `-L` flag to specify the library directory:

```
g++ module.o main.o -o my_program -lmy_lib -L/path/to/my/lib
```

The `-lmy_lib` flag links to the `libmy_lib.so` (dynamic) or `libmy_lib.a` (static) file in the specified directory.

2.11.3 Preprocessor, Compiler, Linker: Simplified Workflow

For small projects with few dependencies, a single command can handle preprocessing, compilation, and linking:

```
g++ mod1.cpp mod2.cpp main.cpp -I/path/to/include/dir -o my_program
```

Warning: Compiler Behavior

Different compilers (e.g., GCC, Clang) may produce varying behaviors, warnings, or errors. For an example of such differences, see this comparison on [GodBolt](#).

2.11.4 Loader

The **loader** is responsible for preparing the executable program for execution:

- Allocates memory for code and data sections.
- Resolves addresses for dynamically linked libraries.
- Starts program execution.

Running an Executable

```
./my_program
```

Dynamic Libraries and `LD_LIBRARY_PATH`

When linking against external dynamic libraries, the loader uses the environment variable `LD_LIBRARY_PATH` to locate them. Ensure the required library paths are included:

```
export LD_LIBRARY_PATH+=:/path/to/my/lib  
./my_program
```

3

Object-Oriented Programming

3.1 Introduction

What is OOP

Object-Oriented Programming (OOP) is a programming paradigm that emphasizes the use of objects to represent real-world entities and concepts.

Definition: *OOP in C++*

C++ is not (only) an OOP language. It is a **multi-paradigm programming language** that supports procedural, object-oriented, and generic programming. C++ allows developers to combine these paradigms effectively for various programming tasks.

Key Principles of OOP

OOP is based on several key principles:

- **Encapsulation:** bundles data (attributes) and the functions (methods) that operate on the data into a single unit called an *object*. This promotes data hiding and reduces code complexity.
- **Inheritance:** allows you to create new classes (derived or child classes) based on existing classes (base or parent classes). It enables code reuse and the creation of class hierarchies.
- **Polymorphism:** enables objects from different derived classes to be handled through a common base class. This facilitates code reuse and provides a way to handle objects abstractly.

RAII (Resource Acquisition Is Initialization)

RAII is a C++ idiom that binds resource management to an object's lifetime. Resources are acquired in the constructor and released in the destructor, ensuring automatic cleanup when the object goes out of scope.

- Eliminates manual resource management and reduces leaks.
- Ensures deterministic resource release, even in case of exceptions.
- Manages memory, file handles, locks, and network connections efficiently.

Advantages of OOP

OOP offers numerous advantages, including:

- **Modularity:** OOP encourages the division of a complex system into smaller, manageable objects, promoting code modularity and reusability.
- **Maintenance:** Objects are self-contained, making it easier to maintain and update specific parts of the code without affecting other parts.
- **Flexibility:** Inheritance and polymorphism provide flexibility, allowing you to extend and modify the behavior of classes without altering their existing code.
- **Readability:** OOP promotes code readability by organizing data and functions related to a specific object within a class.

3.2 Classes and Objects in C++

3.2.1 Classes and Members

In C++, a class encapsulates `data` (member variables) and `behavior` (member functions), providing a blueprint for creating objects. This encapsulation ensures that data and related functionality are bundled together, promoting modularity and reusability.

The following example defines a simple `Car` class that models a real-world vehicle. It includes attributes such as the manufacturer, model, and year, along with a method to start the engine.

```
1 class Car {  
2 public:  
3     std::string manufacturer;  
4     std::string model;  
5     unsigned int year;  
6  
7     void start_engine() {  
8         std::cout << "Engine started!" << std::endl;  
9     }  
10};
```

Creating and Using Objects

Once a class is defined, objects can be instantiated to represent real entities. You can create an instance of the `Car` class and interact with its members using the dot (`.`) operator.

```
11 Car my_car;           // Creating an object of class Car.  
12 my_car.manufacturer = "BMW";  
13 my_car.model = "X5";  
14 my_car.year = 2024;  
15 my_car.start_engine(); // Invoking a method.
```

C++ allows objects to be managed dynamically using pointers. This provides flexibility, especially when working with objects whose lifetimes extend beyond the current scope. The arrow operator (`->`) is used to access members of a pointer to an object.

```
16 Car* my_car_ptr = new Car{};  
17 my_car_ptr->manufacturer = "Alfa Romeo";  
18 my_car_ptr->model = "Giulietta";  
19 my_car_ptr->year = 2010;  
20 my_car_ptr->start_engine(); // Invoking a method.  
21 delete my_car_ptr;        // Releasing allocated memory.
```

Member Variables and Functions

A class consists of two main components:

- **Member Variables:** Store the state of an object. They are also known as attributes or fields.
- **Member Functions:** Define the behavior of an object. They are also known as methods.

In the previous example, the member variables (`manufacturer`, `model`, `engine_running`) store the car's state, while member functions (`start_engine`, `is_running`) define its behavior.

Static Members

Static members are shared by all instances of a class. They are declared with the `static` keyword and accessed using the class name rather than an object. Consider the `Circle` class below, which uses a static constant for `PI`:

```
1 class Circle {
2 public:
3     static const double PI = 3.14159265359;    // Static member.
4     double radius;                           // Non-static member.
5
6     double calculate_area() {
7         return PI * radius * radius;
8     }
9
10    static void print_shape_name() {
11        std::cout << "This is a circle." << std::endl;
12    }
13 }
```

Static members can be accessed using the class name directly, while non-static members require an instance of the object:

```
14 Circle circle;
15 circle.radius = 5.0;
16 const double area = circle.calculate_area(); // non-static member.
17 const double pi_value = Circle::PI;           // static member.
18 Circle::print_shape_name();
```

Const Members

The keyword `const` can be applied to member variables and member functions to indicate immutability. In the example below, the constructor initializes a constant member, and the member function `print_value` is declared as `const` to signal that it does not modify the object.

```
1 class MyClass {
2 public:
3     MyClass(int x) : value(x) {}
4
5     void print_value() const {
6         // value *= 2; // Illegal: cannot modify a const member.
7         std::cout << "Const version: " << value << std::endl;
8     }
9
10    const int value;
11 }
```

Tip: `mutable` Keyword

If a `const` member function needs to modify a member variable, that variable can be declared as `mutable` (although this should be used sparingly).

```
mutable int var = 0; // Can be modified in a const member function
```

③ Example: *Const and Non-Const Member Functions*

Here is an example illustrating both const and non-const versions of a member function:

```
1 class Counter {
2 public:
3     Counter() : count(0) {}
4
5     // Non-const version can modify the object
6     void increment() {
7         count++;
8     }
9
10    // Const version can only read the object
11    int get_count() const {
12        return count;
13    }
14
15 private:
16     int count;
17 };
18
19 Counter c1;      // Non-const object
20 const Counter c2; // Const object
21
22 c1.get_count();  // OK      non-const obj can call const func
23 c1.increment();  // OK      non-const obj can call non-const func
24
25 c2.get_count();  // OK      const obj can call const func
26 c2.increment();  // Error   const obj cannot call non-const func
```

The `this` Pointer

The `this` pointer is automatically passed to non-static member functions and points to the object invoking the function. It is useful for resolving ambiguity between member variables and parameters.

```
1 class MyClass {
2 public:
3     int x;
4
5     void print_x() const {
6         std::cout << "Value of x: " << this->x << std::endl;
7     }
8 };
```

3.2.2 Constructors and Destructors

Constructors are special member functions that initialize objects when they are created. They share the same name as the class and may accept arguments to initialize member variables.

Default Constructor

A default constructor takes no arguments. If a class has no explicitly defined constructor, C++ generates a default constructor automatically.

```

1 class MyClass {
2 public:
3     // Default constructor.
4     MyClass() = default;
5
6     std::string name;
7     unsigned int length;
8 }
9
10 MyClass obj;      // Default initialization.
11 MyClass obj2{}; // Uniform initialization (preferred).
12 MyClass* ptr = new MyClass(); // Calls the default constructor.

```

Observation: Default Constructor and Initialization

If you provide any custom constructors, C++ will not generate a default constructor unless explicitly defined.

Default initialization of primitive types (e.g., `int`, `double`) sets them to zero, while objects may be initialized by their own default constructors.

Parameterized Constructor

A parameterized constructor allows objects to be initialized with specific values.

```

1 class Student {
2 public:
3     Student(std::string name, unsigned int age) {
4         this->name = name;
5         this->age = age;
6     }
7
8     std::string name;
9     unsigned int age;
10 };
11
12 Student student1("Alice", 20); // initialization using a constructor.
13 Student student2{"Bob", 23}; // Uniform initialization.

```

A good practice is to use an **initializer list** to initialize member variables in the constructor:

```

1 class Rect {
2 public:
3     Rect(double length, double width) : length(length), width(width) {
4         // Constructor body (if needed).
5     }
6
7     double calculate_area() const {
8         return length * width;
9     }
10
11     double length;
12     double width;
13 };
14
15 Rect rect{5.0, 3.0}; // initialization using an initializer list.
16 const double area = rect.calculate_area();

```

When using an initializer list, the member variables **must be initialized in the order they are declared** in the class definition.

 **Tip: Initializer List**

Using an initializer list avoids redundant assignments and improves efficiency.

Copy Constructor and Copy Assignment

A copy constructor creates a new object as a copy of an existing one, while the copy assignment operator assigns the contents of one object to another existing object.

```
class Book {
public:
    Book(std::string title, std::string author)
        : title(title), author(author) {}

    // Copy constructor.
    Book(const Book& other)
        : title(other.title), author(other.author) {}

    // Copy assignment operator.
    Book& operator=(const Book& other) {
        if (this != &other) {
            title = other.title;
            author = other.author;
        }
        return *this;
    }

    void display_info() const {
        std::cout << "Title:" << title << ", Author:" << author << "\n";
    }

    std::string title;
    std::string author;
};

Book book1{"The Catcher in the Rye", "J.D. Salinger"};
Book book2 = book1; // Copy constructor.
Book book3{"Marcovaldo", "I. Calvino"};
book3 = book1; // Copy assignment operator.
```

 **Observation: Pass and Return by Value**

When passing or returning objects by value, the copy constructor is invoked to create a copy of the object.

```
1 void some_function(Student s) {
2     // Calls the copy constructor when s is passed.
3 }
4
5 Student create_student() {
6     Student s{"Bob", 22};
7     return s; // Calls the copy constructor when s is returned.
8 }
```

Destructor

A destructor is a special member function that cleans up resources when an object is destroyed. It has the same name as the class but is prefixed with `~`. Destructors are automatically called when:

- An object goes out of scope
- The `delete` operator is used on a pointer to the object
- The program ends and static/global objects are destroyed

```
1 class FileHandler {
2 public:
3     FileHandler(std::string filename) : filename(filename) {
4         file.open(filename);
5     }
6
7     ~FileHandler() {
8         if (file.is_open()) {
9             file.close();
10        }
11    }
12
13    std::string filename;
14    std::ofstream file;
15 };
16
17 { // The object is automatically destroyed when it goes out of scope.
18     FileHandler file{"data.txt"};
19 } // Destructor is called here, ensuring proper resource cleanup.
```

Observation: Rule of Three

Copy constructor, copy assignment operator and destructor are part of the **Rule of Three** in C++, which states that if a class defines any of the following:

- Destructor
- Copy constructor
- Copy assignment operator

then it should provide all three to ensure proper memory management.

Compiler-Generated Special Member Functions

Table 3.1: The compiler automatically generates special member functions when not explicitly defined.
[howardhinnant]

user declares	default constructor	destructor	copy constructor	copy assignment	move constructor	move assignment
Nothing	defaulted	defaulted	defaulted	defaulted	defaulted	defaulted
Any constructor	not declared	defaulted	defaulted	defaulted	defaulted	defaulted
default constructor	user declared	defaulted	defaulted	defaulted	defaulted	defaulted
destructor	defaulted	user declared	defaulted	defaulted	not declared	not declared
copy constructor	not declared	defaulted	user declared	defaulted	not declared	not declared
copy assignment	defaulted	defaulted	defaulted	user declared	not declared	not declared
move constructor	not declared	defaulted	deleted	deleted	user declared	not declared
move assignment	defaulted	defaulted	deleted	deleted	not declared	user declared

3.2.3 The `inline` Directive

The `inline` keyword in C++ suggests that the compiler replaces a function call with the actual function code at the call site. This can reduce function call overhead and improve performance for small, frequently used functions.

Syntax

```
// Inline function definition.  
inline int add(int a, int b) {  
    return a + b;  
}  
  
// Function call.  
const int result = add(5, 7);  
std::cout << "Result: " << result << std::endl;
```

Inline functions have the following characteristics:

- **Limited to small functions:** Inlining is beneficial for short functions but can lead to code bloat for larger functions.
- **Compiler optimization:** The `inline` keyword is only a suggestion; the compiler may choose whether to inline a function based on optimization settings.
- **Usage in header files:** Inline functions should be defined in header files to avoid multiple definition errors, but proper use of header guards is essential.
- **Impact on readability:** Excessive inlining can make debugging harder and lead to poor code organization.

Advantages

- Reduces function call overhead.
- Eliminates stack frame creation for simple functions.
- Helps avoid multiple definitions across translation units.
- Useful for small, performance-critical functions.

Considerations

- Excessive use can increase executable size.
- Compiler may ignore the `inline` directive if deemed inefficient.
- Debugging inlined functions is more difficult.
- Should be used cautiously to maintain code clarity.

💡 Tip: When to Use `inline`

The `inline` directive is most effective when:

- Use for very small functions, such as one-liners.
- Use when function call overhead is significant.
- Avoid inlining large or recursive functions.
- Ensure proper placement in header files to prevent linkage issues.

3.2.4 In-Class and Out-of-Class Definitions

In C++, member functions of a class can be defined either inside the class declaration (in-class, implicitly `inline`) or separately in a source file (out-of-class). Each approach has its advantages and trade-offs.

In-Class Definition (Implicitly Inline)

A member function defined inside the class is automatically considered `inline`. This is common for short functions that are one-liners or concise operations.

```
// my_class.hpp
class MyClass {
public:
    // Inline function defined inside the class.
    int add(int a, int b) {
        return a + b;
    }
};
```

The compiler treats this as if it were explicitly marked `inline` when defined outside the class.

```
// my_class.hpp
class MyClass {
public:
    int add(int a, int b);
};

// Inline keyword is implicit when defined inside the class.
inline int MyClass::add(int a, int b) {
    return a + b;
}
```

Advantages

- Improves readability for simple functions.
- Allows potential inlining by the compiler, reducing function call overhead.

Disadvantages

- May lead to code bloat with large functions.
- Changes to the function force recompilation of all translation units, including header files.

Out-of-Class Definition (Separation of Interface and Implementation)

Member functions can also be declared in the class but defined separately in a source file (`.cpp`). This is typically done for larger functions or when separating interface from implementation.

```
// my_class.hpp
class MyClass {
public:
    int add(int a, int b); // Function declaration.
};
```

```
// my_class.cpp
#include "my_class.hpp"

int MyClass::add(int a, int b) {
    return a + b;
}
```

Advantages

- Improves code organization.
- Changes to function implementation do not require recompilation of all translation units.

Disadvantages

- Slightly more verbose.
- Requires managing separate `.cpp` files.

💡 Tip: Best Practices for Function Definitions

1. Use **in-class definitions** for very short functions (e.g., accessors, mutators) where inlining might improve performance.
2. Use **out-of-class definitions** for more complex functions to keep headers clean and separate interface from implementation.
3. Balance readability, maintainability, and performance when deciding where to define functions.

In practice, a combination of both approaches is used to maintain structured and efficient code.

3.2.5 Encapsulation and Access Control

Encapsulation is a key principle in object-oriented programming (OOP) that bundles data (attributes) and methods (functions) into a single unit: an object.

It restricts direct access to an object's internal state, ensuring data integrity and security.

Encapsulation in C++

Encapsulation allows defining a clear interface while hiding implementation details. Data members should generally be private, with public methods providing controlled access.

💡 Example: Encapsulation Example

```
class BankAccount {
public:
    BankAccount(std::string account_holder, double balance)
        : account_holder(account_holder), balance(balance) {}

    void deposit(double amount) {
        balance += amount;
    }

    double get_balance() const {
        return balance;
    }

private:
    std::string account_holder;
    double balance;
};
```

Access Specifiers in C++

C++ provides three access specifiers to control member visibility:

- `public` : Accessible from any part of the program (forms the class's public interface).
- `private` : Only accessible within the class itself (internal implementation).
- `protected` : Similar to `private`, but also accessible in derived classes (used in inheritance).

```

class MyClass {
public:
    int public_var;      // Accessible from anywhere.
    void public_func() { /* ... */ }

private:
    int private_var;    // Accessible only within this class.
    void private_func() { /* ... */ }
};

```

3.2.6 Class vs. Struct

In C++, both `class` and `struct` can be used to define objects with member variables and functions.

The primary difference between them lies in their **default access specifiers**:

- In a `class`, members are **private** by default.
- In a `struct`, members are **public** by default.

```

// Class vs. Struct Example
class MyClass {
    int x; // Private by default.
public:
    int y; // Explicitly public.
};

struct MyStruct {
    int x; // Public by default.
private:
    int y; // Explicitly private.
};

```

When to Use Class vs. Struct

While functionally equivalent, structs and classes serve different purposes in common C++ conventions.

Use `class` when:

- Encapsulating data with private members
- Defining objects with complex behavior
- Implementing OOP concepts (e.g. inheritance)
- Maintaining strong encapsulation

Use `struct` when:

- Grouping related public variables
- Using simple data structures (e.g. point, color)
- Interacting with C-style structs or APIs
- Defining POD (Plain Old Data) types

3.2.7 Getter and Setter Methods

Getter and setter methods (also known as **accessors and mutators**) provide controlled access to private member variables.

- **Getters:** Allow reading private data.
- **Setters:** Allow modifying private data with validation or restrictions.

Example: *Getter and Setter*

```
1 class TemperatureSensor {
2 public:
3     double get_temperature() const {
4         return temperature;
5     }
6
7     void set_temperature(double new_temperature) {
8         if (new_temperature >= -50.0 && new_temperature <= 150.0) {
9             temperature = new_temperature;
10        } else {
11            std::cout << "Invalid temperature value!" << std::endl;
12        }
13    }
14
15 private:
16     double temperature;
17 };
```

Using getters and setters enforces encapsulation while allowing controlled data access.

3.2.8 Friend Classes

A **friend** class allows another class to access its private and protected members. This is useful when two classes need close interaction but maintaining strict encapsulation is impractical.

Let's consider a **Circle** class with a private member **radius**.

```
// circle.hpp
class Circle {
public:
    friend class Cylinder;
    // Cylinder class has access to private members of Circle.

    Circle(double r) : radius(r) {}

    double get_area() const {
        return 3.14159265359 * radius * radius;
    }

private:
    double radius;
};
```

Let's consider now a **Cylinder** class that requires access to the private member **radius** of the **Circle** class:

```

// cylinder.hpp
class Cylinder {
public:
    Cylinder(const Circle &circle, double height)
        : circle(circle), height(height) {}

    double get_volume() const {
        // Accessing the private member 'radius' of Circle.
        return circle.radius * circle.radius * height;
    }

private:
    double height;
    const Circle circle;
};

Circle circle{1.0};
Cylinder cylinder{circle, 0.5};
const double volume = cylinder.get_volume();

```

 **Tip: Best Practices for Friend Classes**

- **Minimize usage:** Prefer encapsulation unless two classes require deep interaction.
- **Use only when necessary:** Friend classes break encapsulation, so they should be well justified.
- **Maintain modularity:** Avoid excessive interdependencies between classes.

Friend classes allow controlled access to private members but should be used sparingly to maintain encapsulation principles.

3.2.9 Operator Overloading

Operator overloading in C++ allows defining custom behaviors for operators when used with user-defined classes. This makes objects behave more naturally and improves code readability.

Category	Operators
Arithmetic	+ , - , * , / , %
Comparison	== , != , < , > , ≤ , ≥ , ⇔ (C++20)
Assignment	= , += , -= , *=
Increment/Decrement	++ , --
Stream	<< (output), >> (input)
Subscript	[] (array-like behavior)
Function Call	() (used in functor classes)
Pointer	→ , * (smart pointers, iterators)

Table 3.2: Commonly Overloaded Operators in C++

By overloading operators you can extend their functionality for custom types. Instead of writing verbose function calls, operator overloading lets you use familiar syntax for objects.

⌚ Example: *Operator Overloading*

```
class Complex {
public:
    Complex(double r, double i) : real(r), imag(i) {}

    // Overloading the + operator.
    Complex operator+(const Complex& other) {
        return Complex(real + other.real, imag + other.imag);
    }

    void print() const {
        std::cout << real << " + " << imag << "i" << std::endl;
    }

private:
    double real, imag;
};

Complex a{2.0, 3.0};
Complex b{1.0, 2.0};
Complex c = a + b; // Using the overloaded '+' operator.
c.print();
```

Operators can be overloaded as **member functions** or **non-member functions**.

Overloading as a Member Function

If the left operand is an object of the class, the operator can be overloaded as a member function.

```
// Overloading as a Member Function
class Vector {
public:
    Vector(int x, int y) : x(x), y(y) {}

    Vector operator+(const Vector& other) {
        return Vector(x + other.x, y + other.y);
    }

    void print() const {
        std::cout << "(" << x << ", " << y << ")" << std::endl;
    }

private:
    int x, y;
};

Vector v1{2, 3}, v2{1, 1};
Vector result = v1 + v2; // Works because + is a member function.
result.print();
```

Overloading as a Non-Member Function

If the left operand is not an object of the class, the operator must be overloaded as a non-member function.

```
// Overloading as a Non-Member Function (Friend Function)
class MyClass {
public:
    MyClass(int v) : value(v) {}

    // Declaring the '<<' operator as a friend function.
    friend std::ostream& operator<<(std::ostream& os, const MyClass& obj
        );

private:
    int value;
};

// Overloading the '<<' operator as a non-member function.
std::ostream& operator<<(std::ostream& os, const MyClass& obj) {
    os << obj.value;
    return os;
}

MyClass obj{42};
std::cout << obj << std::endl; // Calls the overloaded operator.
```

💡 Tip: *Operator Overloading Guidelines*

1. Operators that cannot be overloaded

Some operators like `::`, `.*`, and `? :` cannot be overloaded.

2. Preserve the operator's meaning

Overloading should make sense in the class's context (e.g., `+` for addition).

3. Respect operator precedence

Overloaded operators should follow the same precedence as their built-in counterparts.

4. Avoid excessive overloading

Overloading too many operators can make code hard to maintain.

Operator overloading enhances code readability when used correctly, making custom types feel like native types.

4

Standard Template Library

The Standard Template Library (STL) is a collection of C++ template classes that provide general-purpose data structures and algorithms, including `std::vector`, `std::list`, `std::queue`, and `std::stack`. The STL consists of four main components:

- **Algorithms:** A collection of functions designed to operate on ranges of elements.
- **Containers:** Objects that store collections of other objects.
- **Function Objects:** Components that allow the creation of callable objects.
- **Iterators:** Objects that enable traversal of a container.

4.1 Containers

Containers are objects that store data. The STL provides several container classes, each supporting different operations. STL containers are categorized as follows:

- **Sequence Containers:** These containers maintain ordered collections of elements. The main sequence containers are `std::vector`, `std::list`, and `std::deque`.
- **Associative Containers:** These containers store elements in sorted order and support fast searching. The main associative containers are `std::set`, `std::multiset`, `std::map`, and `std::multimap`.
- **Container Adaptors:** These provide restricted access to elements by adapting existing containers. The primary container adaptors are `exttstd::stack`, `exttstd::queue`, and `exttstd::priority_queue`.
- **Special Containers:** These containers provide specialized functionality. Examples include `std::byte`, `std::pair`, `std::tuple`, `std::variant`, `std::optional` and `std::any`.

4.1.1 Sequence Containers

Sequence containers maintain ordered collections of elements, with element positions independent of their values. In `std::vector` and `std::array`, elements are stored contiguously in memory and can be accessed using an index `[]`.

Examples of sequence containers include `std::vector<T>`, `std::array<T, N>`, `std::deque<T>`, and `std::list<T>`.

③ Example:

```
1 #include <iostream>
2 #include <vector>
3
4 int main() {
5     std::vector<int> v = {1, 2, 3, 4, 5};
6     for (int num : v) {
7         std::cout << num << " ";
8     }
9     return 0;
10 }
```

4.1.2 Associative Containers

Associative containers store elements in sorted order and support fast searching. The main associative containers are `std::set`, `std::multiset`, `std::map`, and `std::multimap`.

- `std::set` is a container that stores unique elements in sorted order. Here, **key** and **value** are the same and thus used interchangeably.
- `std::multiset` is a container that stores multiple elements in sorted order.
- `std::map` is a container that stores key-value pairs in sorted order.
- `std::multimap` is a container that stores multiple key-value pairs in sorted order.

③ Example:

```
1 #include <iostream>
2
3 int main() {
4     std::map<std::string, int> m;
5     m["one"] = 1;
6     m["two"] = 2;
7     m["three"] = 3;
8
9     for (const auto& [key, value] : m) {
10         std::cout << key << " => " << value << std::endl;
11     }
12     return 0;
13 }
```

They can be further divided in **ordered** and **unordered** associative containers. The former maintain elements in sorted order, while the latter do not. For the former, an ordering relation is required for the elements, which can be provided by a comparison function or by a comparison operator. Moreover, keys can be accessed read-only, but not modified. For the latter, a hashing function is required for the elements, which can be provided by a hash function or by a hash operator. Keys can be accessed and modified.

4.1.3 Container Adaptors

Container adaptors provide restricted access to elements by adapting existing containers. The primary container adaptors are `std::stack`, `std::queue`, and `std::priority_queue`.

- `std::stack` is a container that provides a LIFO (Last In, First Out) data structure.
- `std::queue` is a container that provides a FIFO (First In, First Out) data structure.
- `std::priority_queue` is a container that provides a priority queue data structure.

?

Example:

```

1 #include <iostream>
2 #include <queue>
3
4 int main() {
5     std::queue<int> q;
6     q.push(1);
7     q.push(2);
8     q.push(3);
9
10    while (!q.empty()) {
11        std::cout << q.front() << " ";
12        q.pop();
13    }
14    return 0;
15 }
```

4.1.4 Special Containers

Special containers provide specialized functionality. Examples include `std::byte`, `std::pair`, `std::tuple`, `std::variant`, `std::optional`, and `std::any`.

- `std::byte` is a container that stores byte values.
- `std::pair` is a container that stores a pair of elements.
- `std::tuple` is a container that stores a tuple of elements.
- `std::variant` is a container that stores a variant of elements.
- `std::optional` is a container that stores an optional value.
- `std::any` is a container that stores any type of value.

?

Example:

```

1 #include <iostream>
2 #include <tuple>
3
4 int main() {
5     std::tuple<int, float, std::string> t(1
6         , 3.14, "Hello");
7     std::cout << std::get<0>(t) << " ";
8     std::cout << std::get<1>(t) << " ";
9
10    return 0;
11 }
```

For further examples see [here](#).

4.1.5 Iterators

Iterators are a generalization of **pointers** that allow a C++ program to work with different data structures (for example, containers and ranges (since C++20)) in a uniform manner. The iterator library provides definitions for iterators, as well as iterator traits, adaptors, and utility functions.

Definition:

An iterator is any object that allows iterating over a succession of elements, typically stored in a standard container. It can be dereferenced with the `*` operator, returning an element of the range, and incremented (moving to the next element) with the `++` operator.

Iterator category	Operations and storage requirement						
	write	read	increment		decrement	random access	contiguous storage
			without multiple passes	with multiple passes			
<i>Iterator</i>			Required				
<i>OutputIterator</i>	Required		Required				
<i>InputIterator</i> (mutable if supports write operation)		Required	Required				
<i>ForwardIterator</i> (also satisfies <i>InputIterator</i>)		Required	Required	Required			
<i>BidirectionalIterator</i> (also satisfies <i>ForwardIterator</i>)		Required	Required	Required	Required		
<i>RandomAccessIterator</i> (also satisfies <i>BidirectionalIterator</i>)		Required	Required	Required	Required	Required	
<i>ContiguousIterator</i> ^[1] (also satisfies <i>RandomAccessIterator</i>)		Required	Required	Required	Required	Required	Required

Figure 4.1: Iterators

Key Concepts of Iterators

- Abstraction:** Iterators provide a way to access elements of a container without exposing its internal structure.
- Uniform Interface:** They offer a consistent interface for traversing different types of containers (e.g., arrays, vectors, lists, etc.).
- Categories:** Iterators are categorized based on their functionality. The main categories are:
 - Input Iterators:** Can read elements in a sequence (forward-only, single-pass).
 - Output Iterators:** Can write elements in a sequence (forward-only, single-pass).
 - Forward Iterators:** Can read and write elements in a sequence (forward-only, multi-pass).
 - Bidirectional Iterators:** Can move both forward and backward in a sequence.
 - Random Access Iterators:** Can access any element in constant time (like pointers).

Common Iterator Operations

- `*iter`: Dereference the iterator to access the element it points to.
- `iter->member`: Access a member of the element the iterator points to.
- `+iter` / `iter++`: Move the iterator to the next element.

- `--iter` / `iter--` : Move the iterator to the previous element (for bidirectional and random access iterators).
- `iter1 == iter2` / `iter1 != iter2` : Compare two iterators for equality.
- `iter + n` / `iter - n` : Move the iterator by `n` positions (for random access iterators).

Iterator Types in C++ Standard Library

- `begin()` and `end()` :
 - `begin()` returns an iterator to the first element of the container.
 - `end()` returns an iterator to one past the last element (used as a sentinel value).
- `const_iterator` : A read-only iterator that cannot modify the elements of the container.
- `reverse_iterator` : Iterates over the container in reverse order.

Example:

```

1 #include <iostream>
2 #include <vector>
3
4 int main() {
5     std::vector<int> vec = {1, 2, 3, 4, 5};
6
7     // Using iterators to traverse the vector
8     for (auto it = vec.begin(); it != vec.end(); ++it) {
9         std::cout << *it << " "; // Dereference the iterator to
10        access the element
11    }
12    std::cout << std::endl;
13
14    // Using range-based for loop (internally uses iterators)
15    for (int val : vec) {
16        std::cout << val << " ";
17    }
18    std::cout << std::endl;
19
20    return 0;
}

```

Iterator Categories in Practice

- **Random Access Iterators**: Supported by `std::vector`, `std::array`, and `std::deque`.
- **Bidirectional Iterators**: Supported by `std::list` and `std::set`.
- **Forward Iterators**: Supported by `std::forward_list`.
- **Input/Output Iterators**: Used in specific scenarios like reading from or writing to streams.

Custom Iterators

You can also define custom iterators for your own container classes by implementing the required operations (e.g., `+`, `*`, `==`, etc.).

4.1.6 `size_type` and `size_t`

They are used to represent the size of a container. `size_type` is a type defined by the container class, while `size_t` is a type defined by the C++ standard library. They are guaranteed to be unsigned integral types, but their sizes may vary depending on the platform.

Example:

```
1 #include <iostream>
2 #include <vector>
3
4 int main() {
5     std::vector<int> vec = {1, 2, 3, 4, 5};
6     std::vector<int>::size_type vec_size = vec.size();
7     std::size_t vec_size_t = vec.size();
8
9     std::cout << "size_type: " << vec_size << std::endl;
10    std::cout << "size_t: " << vec_size_t << std::endl;
11
12    return 0;
13 }
```

4.2 Algorithms

The STL provides an extensive set of algorithms to operate on containers, or more precisely on **ranges**.

Warning:

C++20 has revised the concept of range and provides a new set of algorithms in the namespace `std::ranges`, with the same name as the old ones, but simpler to use and sometimes more powerful.

Why use STL Algorithms?

With standardized algorithms:

- You are more uniform with respect to different container types.
- The algorithm of the standard library may do certain optimizations if the contained elements have some characteristics.
- You have a parallel version for free...

4.2.1 Non-modifying Algorithms

Non-modifying algorithms do not change the contents of the container (the value in the range). They are used to perform operations like searching, counting, and comparing elements.

③ Example:

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4
5 int main() {
6     std::vector<int> vec = {1, 2, 3, 4, 5};
7
8     // Find the first element greater than 3
9     auto it = std::find_if(vec.begin(), vec.end(), [] (int x) {
10         return x > 3; });
11
12     if (it != vec.end()) {
13         std::cout << "First element greater than 3: " << *it << std
14             ::endl;
15     }
16
17     // Check if all elements are positive
18     bool all_positive = std::all_of(vec.begin(), vec.end(), [] (int
19         x) { return x > 0; });
20
21     if (all_positive) {
22         std::cout << "All elements are positive" << std::endl;
23     }
24 }
```

4.2.2 Modifying Algorithms

Modifying algorithms change the contents of the container. They are used to perform operations like sorting, removing elements, and transforming elements.

③ Example:

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4
5 int main() {
6     std::vector<int> vec = {5, 2, 3, 4, 1};
7
8     // Sort the elements in ascending order
9     std::sort(vec.begin(), vec.end());
10
11    // Remove the first element
12    vec.erase(vec.begin());
13
14    // Transform each element to its square
15    std::transform(vec.begin(), vec.end(), vec.begin(), [] (int x) {
16        return x * x; });
17
18    for (int val : vec) {
19        std::cout << val << " ";
20    }
21    std::cout << std::endl;
22
23    return 0;
24 }
```

4.2.3 Inserters

Inserters are used to insert elements into a container. The main inserters are:

- `std::back_inserter`: Inserts elements at the end of a container.
- `std::front_inserter`: Inserts elements at the beginning of a container.
- `std::inserter`: Inserts elements at a specified position in a container.

② Example:

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4
5 int main() {
6     std::vector<int> vec1 = {1, 2, 3};
7     std::vector<int> vec2 = {4, 5, 6};
8
9     // Insert elements from vec2 to vec1
10    std::copy(vec2.begin(), vec2.end(), std::back_inserter(vec1));
11
12    for (int val : vec1) {
13        std::cout << val << " ";
14    }
15    std::cout << std::endl;
16
17    return 0;
18 }
```

4.2.4 Sorting Algorithms

Sorting algorithms are used to arrange elements in a specific order. They operate on a range to order it according to an ordering relation.

② Example:

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4
5 int main() {
6     std::vector<int> vec = {5, 2, 3, 4, 1};
7
8     // Sort the elements in ascending order
9     std::sort(vec.begin(), vec.end());
10
11    for (int val : vec) {
12        std::cout << val << " ";
13    }
14    std::cout << std::endl;
15
16    return 0;
17 }
```

4.2.5 Min and Max

They are used to find the minimum and maximum elements in a range. The functions `std::min_element` and `std::max_element` return an iterator to the minimum and maximum elements, respectively.

② Example:

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4
5 int main() {
6     std::vector<int> vec = {5, 2, 3, 4, 1};
7
8     // Find the minimum and maximum elements
9     auto min_it = std::min_element(vec.begin(), vec.end());
10    auto max_it = std::max_element(vec.begin(), vec.end());
11
12    std::cout << "Min element: " << *min_it << std::endl;
13    std::cout << "Max element: " << *max_it << std::endl;
14
15    return 0;
16 }
```

4.2.6 Numeric Algorithms

Numeric algorithms are used to perform numeric operations on a range of elements. The main numeric algorithms are:

- `std::accumulate` : Computes the sum of elements in a range.
- `std::inner_product` : Computes the inner product of two ranges.
- `std::partial_sum` : Computes the partial sum of elements in a range.
- `std::adjacent_difference` : Computes the differences between adjacent elements in a range.

② Example:

```
1 #include <iostream>
2 #include <vector>
3 #include <numeric>
4
5 int main() {
6     std::vector<int> vec = {1, 2, 3, 4, 5};
7
8     // Compute the sum of elements
9     int sum = std::accumulate(vec.begin(), vec.end(), 0);
10
11    // Compute the partial sum of elements
12    std::vector<int> partial_sum(vec.size());
13    std::partial_sum(vec.begin(), vec.end(), partial_sum.begin());
14
15    for (int val : partial_sum) {
16        std::cout << val << " ";
17    }
18    std::cout << std::endl;
19
20    return 0;
21 }
```

4.3 STL evolution

The STL has evolved over time, with new features and improvements introduced in each C++ standard. Here are some key changes in the STL from C++11 to C++20:

- **C++11:** Introduced move semantics, lambda expressions, and range-based for loops.
- **C++14:** Added generic lambdas, variable templates, and binary literals.
- **C++17:** Introduced parallel algorithms, `std::optional`, `std::variant`, and `std::any`.
- **C++20:** Added concepts, ranges, coroutines, and `std::span`.

For loop evolution

The range-based for loop was introduced in C++11 to simplify iteration over containers. It allows you to iterate over a range of elements without using iterators explicitly.

Example:

```
1 #include <iostream>
2 #include <vector>
3
4 int main() {
5     std::vector<int> vec = {1, 2, 3, 4, 5};
6
7     // Using range-based for loop
8     for (int val : vec) {
9         std::cout << val << " ";
10    }
11    std::cout << std::endl;
12
13    return 0;
14 }
```

4.4 Smart Pointers

Definition: RAII

Resource Acquisition Is Initialization plays a significant role in C++. It essentially means that an object should be responsible for both the creation and destruction of the resources it owns.

Example: RAII Compliant

```
1 std::array<double, 10> p; // var p creates and destroys the array
```

In C++, **smart pointers** are important tools to implement RAII.

Types of pointers:

- **Standard Pointers**: use them only to watch and operate on an object whose lifespan is independent of the one of the pointer.
- **Smart Pointers**: they control the lifespan of an object.
 - `std::unique_ptr`: manages a single object. The owned resource is destroyed when the pointer goes out of scope.
 - `std::shared_ptr`: manages a shared object. The owned resource is destroyed when the **last** shared pointer goes out of scope.
 - `std::weak_ptr`: observes a shared object without owning it. It is used to break circular references between `std::shared_ptr`s.

Example:

```
1 #include <iostream>
2
3 class MyClass {
4 public:
5     set_polygon(std::unique_ptr<Polygon> p) {
6         polygon = std::move(p);
7     }
8 private:
9     std::unique_ptr<Polygon> polygon;
10};
11
12 std::unique_ptr<Polygon> create_polygon() {
13     switch(t){
14         case "Triangle":
15             return std::make_unique<Triangle>();
16         case "Rectangle":
17             return std::make_unique<Rectangle>();
18         default:
19             return nullptr;
20     }
21 }
22
23 MyClass obj;
24 obj.set_polygon(create_polygon("Triangle"));
```

4.4.1 std::unique_ptr

A `std::unique_ptr` serves as the unique owner of the object of type `T` it refers to. The object is destroyed automatically when the `std::unique_ptr` goes out of scope. It implements the `*` and `→` operators to access the object it owns, so it can be used as standard pointer, but can only be initialized through the constructor.

The default constructor produces a **null pointer**.

⚠ Warning: *Copying*

By definition, unique pointers cannot be **copied**, but their ownership can be transferred using the `std::move` function. The previous owner will be left with a null pointer.

4.4.2 std::shared_ptr

For instance you have several objects that refer to a resource (e.g., a matrix, a shape, ...) that is build dynamically (and maybe is a polymorphic object). You want to keep track of all the references in such a way that when (and only when) the last one gets destroyed the resource is also destroyed.

A **Shared Pointer** implements the semantics of *clean it up* when the resource is no longer needed.

❓ Example:

```
1 #include <iostream>
2
3 class MyClass { ... };
4
5 class Preprocessor {
6 public:
7     Preprocessor(std::shared_ptr<Data> &data, ...) : data(data) {}
8 private:
9     std::shared_ptr<Data> data;
10};
11
12 class NumericalSolver {
13 public:
14     NumericalSolver(std::shared_ptr<Data> &data, ...) : data(data)
15     {}
16 private:
17     std::shared_ptr<Data> data;
18 };
19 int main() {
20     std::shared_ptr<Data> data = std::make_shared<Data>();
21     Preprocessor preprocessor(data, ...);
22     preprocessor.preprocess();
23     // shared\_data will still be used by other resources, hence it
24     // cannot be destroyed here.
25     NumericalSolver solver(data, ...);
26     solver.solve();
27     return 0;
28 }
```

How a shared pointer works

The `std::shared_ptr` keeps track of the number of shared references to an object through reference counting. When the reference count reaches zero, the object is automatically deallocated, preventing memory leaks. It implements `*` and `→` operators to access the object it owns, so it can be used as a standard pointer. Moreover, it provides copy constructors and assignment operators.

4.4.3 std::weak_ptr

The `std::weak_ptr` is a smart pointer that holds a non-owning (weak) reference to an object managed by a `std::shared_ptr`. It is used to break circular references between shared pointers. It must also be converted to a `std::shared_ptr` to access the object it refers to.

Example:

```
1 #include <iostream>
2
3 std::shared_ptr<int> shared = std::make_shared<int>(42);
4 std::weak_ptr<int> weak = shared; // Get pointer to data without
                                    // taking ownership.
```

4.5 Move Semantics

Problem: how can I swap two objects in an efficient way?

Before C++11 one could have used a special method or a friend function.

Example:

```
1 #include <iostream>
2
3 void swap_with_move(Matrix& a, Matrix& b) {
4     // swap rows and cols
5
6     double* tmp = a.data; // save the pointer
7     a.data = b.data // copy the pointer
8     b.data = tmp; // copy the pointer
9 }
```

But this way it is not generalizable, since I cannot write a function template `swap_with_move<T>`, because I need to know how data is stored in `T` for each case.

Lvalues

Definition:

An **lvalue** (left value) is an expression that refers to a specific memory location and persists beyond a single expression.

- They have an **identifiable memory location**.
- They can be **modified** (unless they are `const`).

- They can be **bound to lvalue references** (`T&`).

Example:

```

1 int x = 10; // 'x' is an lvalue
2 x = 20;      // Valid: lvalues can appear on the left side of an
               assignment
3
4 int& ref = x; // Valid: lvalues can be bound to lvalue references

```

A `const` lvalue is still an lvalue but cannot be modified:

```

1 const int y = 5;
2 y = 10; // Error: lvalue is read-only

```

Rvalues

Definition:

An **rvalue** (right value) is an expression that does **not persist beyond a single expression** and does not have a specific memory location that can be accessed.

- They are **temporary** and cannot be assigned to.
- They **cannot be bound to lvalue references** (`T&`), but they **can be bound to rvalue references** (`T&&`).

Example:

```

1 int x = 10;
2 int y = x + 5; // 'x + 5' is an rvalue

```

C++11 introduced **rvalue references** to allow functions to take ownership of temporary values:

```

1 void foo(int&& a) { // Accepts only rvalues
2     std::cout << "Rvalue: " << a << std::endl;
3 }
4
5 foo(10); // Valid: '10' is an rvalue
6 int x = 5;
7 foo(x); // Error: 'x' is an lvalue

```

Lvalues vs. Rvalues in Function Overloading

Functions can be overloaded based on lvalue and rvalue references:

```

1 void process(int& x) {
2     std::cout << "Lvalue reference" << std::endl;
3 }
4
5 void process(int&& x) {

```

```

6     std::cout << "Rvalue reference" << std::endl;
7 }
8
9 int main() {
10    int a = 10;
11    process(a); // Calls lvalue version
12    process(20); // Calls rvalue version
13 }
```

4.5.1 How Move Semantics is implemented

The following is the standard signature of a move operation for a class `Matrix`:

```

1 Matrix(Matrix&&); // move constructor
2 Matrix& operator=(Matrix&&); // move assignment operator
```

⌚ Example:

```

1 #include <iostream>
2
3 Matrix(Matrix&& rhs) : data(rhs.data), nr(rhs.nr), nc(rhs.nc) {
4     rhs.data = nullptr;
5     rhs.nr = 0;
6     rhs.nc = 0;
7 }
8
9 Matrix& operator=(Matrix&& rhs) {
10     delete[] this -> data; // release the resource
11     data = rhs.data; // shallow copy
12     // Fix rhs so it is a valid empty matrix
13     rhs.data = nullptr;
14     rhs.nr = 0;
15     rhs.nc = 0;
16 }
17
18 int main() {
19     Matrix foo();
20     // ...
21     Matrix a;
22     a = foo(); // move assignment is called
23     return 0;
24 }
```

⌚ Observation: `std::move`

The `std::move` function is used to convert an lvalue to an rvalue reference, allowing it to be moved instead of copied. It basically casts the value to an rvalue, making it available to be copied.

```

1 template<class T>
2 void swap(T& a, T& b) {
3     T tmp = std::move(a);
```

```

4         a = std::move(b);
5         b = std::move(tmp);
6     }
7 // or simpler
8 std::swap(a, b);

```

The solution is to force the move:

```

1  class Foo{
2  public:
3      Foo(Matrix&& m) : my_m(std::move(m)) {}
4      // ...
5  private:
6      Matrix my_m;
7 }

```

Now, `m` is moved into `my_m`.

All standard containers support move semantic, and all standard algorithms are written so that if the contained type implements move semantics, the creation of unnecessary temporaries can be avoided. All containers also have a `swap()` method that performs swaps intelligently.

4.6 Exceptions

A **function** can be seen as a mapping from input data to output data. The conditions under which the input data is considered valid is called **precondition**, while the conditions under which the output data is considered valid is called **postcondition**. Failure to meet these conditions is considered a **failure or bug**.

An **invariant** is a condition that must always be true during the execution of a program. An object is considered to be in an **inconsistent state** if the invariants are not met.

Definition: *Exception*

An **exception** is an anomalous condition that disrupts the normal flow of a program's execution when left unhandled. It is not the result of incorrect coding but rather arises from challenging or unpredictable circumstances.

4.6.1 Run-time assertions

Example:

```

1 double calculate(double operand1, double operand2) {
2     assert(operand2 != 0.0 && "Division by zero");
3     return operand1 / operand2;
4 }

```

For improved efficiency, all assertions can be disabled by defining the `NDEBUG` macro:

```
1 g++ -DNDEBUG -o program program.cpp
```

4.6.2 Compile-time assertions

Example:

```
1 template <typename T>
2 void process(T value) {
3     static_assert(std::is_integral<T>::value, "T must be an
4         integral type");
5     // Process the value
6 }
```

If the condition is met, the error message is printed to the standard error and compilation will fail.

4.6.3 Exception handling in C++

The basic structure to handle exceptions is:

- Using the `throw` command to indicate that an exception has occurred. You can throw an object containing information about the exception.
- Employing the `try-catch` blocks to catch and handle exceptions. If an exception is not caught, it will propagate up the call stack and might lead to program termination.

Example:

```
1 int divide(int dividend, int divisor) {
2     if (divisor == 0) {
3         throw std::runtime_error("Division by zero");
4     }
5     return dividend / divisor;
6 }
7
8 try{
9     const int result = divide(10, 0);
10    std::cout << "Result: " << result << std::endl;
11 } catch (const std::runtime_error& e) {
12     std::cerr << "Exception: " << e.what() << std::endl;
13 }
```

Standard exceptions examples are:

- `std::exception`: Base class for all standard exceptions.
- `std::runtime_error`: Exception used to report runtime errors.
- `std::logic_error`: Exception used to report errors in the program's logic.
- `std::invalid_argument`: Exception used to report invalid arguments.
- `std::out_of_range`: Exception used to report out-of-range errors.
- `std::bad_alloc`: Exception used to report memory allocation errors.
- `std::bad_cast`: Exception used to report casting errors.

Observation:

In situations where an algorithm's failure is one of its expected outcomes (e.g., the failure of convergence in an iterative method), returning a **status** rather than throwing an exception may be more suitable. Instead of terminating the program, a status variable is used to indicate the outcome, which can be checked by the caller. See also `std::terminate`, `std::abort` and `std::exit`.

Note, also, that **floating-point exceptions** do not result in program failure, instead they produce special numerical values like `Nan` or `Inf`, and the operations continue.

4.7 STL Utilities

4.7.1 I/O Streams

Input/Output (I/O) streams in C++ provide a convenient way to perform input and output operations, allowing you to work with various data sources and destinations, such as files, standard input/output, strings, and more.

The key components are:

- **iostream**: Provides the basic input/output operations. It is used for interacting with the standard input and output streams.
- **ifstream**: This class is used for reading data from files. You can open a file for input and read data from it.

```
1     std::ifstream file("input.txt");
2
3     if (file.is_open()) {
4         std::string line;
5         while (std::getline(file, line)) {
6             std::cout << line << std::endl;
7         }
8         file.close();
9     } else {
10         std::cerr << "Failed to open the file." << std::endl;
11     }
```

- **ofstream**: This class is used for writing data to files. You can open a file for output and write data to it.

```
1     std::ofstream file("output.txt");
2
3     if (file.is_open()) {
4         file << "Hello, World!" << std::endl;
5         file.close();
6     } else {
7         std::cerr << "Failed to open the file." << std::endl;
8     }
```

- **stringstream**: This class is used for reading and writing data to and from strings. It provides a way to work with strings as if they were streams.

```

1     std::stringstream ss;
2     ss << "Hello, World!";
3     std::string str = ss.str();
4     std::cout << str << std::endl;
5
6     // Passing data from a string using std::stringstream.
7     std::string data = "42 3.14 Hello";
8     std::stringstream ss(data);
9     int num;
10    double d;
11    std::string word;
12    ss >> num >> d >> word;

```

4.7.2 Random Numbers

The capability of generating random numbers is essential not only for statistical purposes but also for internet communications. But an algorithm is deterministic. However, several techniques have been developed to generate pseudo-random numbers. They are not really random, but they show a low level of auto-correlation.

To generate them, you need the `<random>` header, and the chosen design is based on two types of objects:

- **Engines**: They are the source of randomness. They are used to generate random numbers.
- **Distributions**: They are used to transform the random numbers generated by the engine into a specific range.

Engines

Random number engines generate pseudo-random numbers using seed data as an entropy source. Several different classes of pseudo-random number generation algorithms are implemented as templates that can be customized.

For simplicity, the library provides predefined engines, such as `std::default_random_engine`, which balances efficiency and quality. There are also non-deterministic engines, like `std::random_device`, which generate non-deterministic random numbers based on hardware data.

Example:

```

1 std::default_random_engine engine; // with default seed
2 std::default_random_engine engine(42); // with seed 42

```

Distributions

Distributions are template classes that implement a call operator () to transform a random sequence into a specific distribution. You need to pass a random engine to the distribution to generate numbers according to the desired distribution. For example:

③ Example:

```
1 std::random_device rd;
2 std::default_random_engine engine(rd());
3 std::uniform_int_distribution<int> dist(1, 6); // [1, 6]
4 int num = dist(engine);
```

4.7.3 Shuffling

In C++ you can shuffle a range of elements using the `std::shuffle` algorithm from the `<algorithm>` header. It rearranges the elements in the range `[first, last)` randomly using a random number generator.

③ Example:

```
1 std::vector<int> vec = {1, 2, 3, 4, 5};
2 std::random_device rd;
3 std::mt19937 engine(rd());
4 std::shuffle(vec.begin(), vec.end(), engine);
```

4.7.4 Sampling

The `std::sample` algorithm from the `<algorithm>` header is used to sample elements from a range. It selects a random subset of elements from the input range and stores them in the output range.

③ Example:

```
1 std::vector<int> vec = {1, 2, 3, 4, 5};
2 std::vector<int> sample(3);
3 std::random_device rd;
4 std::mt19937 engine(rd());
5 std::sample(vec.begin(), vec.end(), sample.begin(), 3, engine);
```

4.7.5 Time Measuring

C++ provides three types of clock:

- `std::chrono::system_clock`: Represents the system-wide real time wall clock.
- `std::chrono::steady_clock`: Represents the monotonic clock that is not affected by system clock adjustments.
- `std::chrono::high_resolution_clock`: Represents the clock with the shortest tick period available on the system.

③ Example:

```
1 auto start = std::chrono::high_resolution_clock::now();
2 // Perform some operation
3 auto end = std::chrono::high_resolution_clock::now();
4 auto duration = std::chrono::duration_cast<std::chrono::
    milliseconds>(end - start);
```

4.7.6 Filesystem Utilities

The C++17 standard introduced the `<filesystem>` library, which provides facilities for performing operations on file systems and their components, such as paths, regular files, and directories.

Key Components

- `std::filesystem::path` : Represents a path in the filesystem.
- `std::filesystem::directory_entry` : Represents a directory entry.
- `std::filesystem::directory_iterator` : Iterates over the contents of a directory.
- `std::filesystem::recursive_directory_iterator` : Recursively iterates over the contents of a directory.

Common Operations

- `std::filesystem::exists` : Checks if a file or directory exists.
- `std::filesystem::create_directory` : Creates a new directory.
- `std::filesystem::remove` : Removes a file or directory.
- `std::filesystem::rename` : Renames a file or directory.
- `std::filesystem::copy` : Copies a file or directory.

② Example:

```
1 #include <iostream>
2 #include <filesystem>
3
4 int main() {
5     std::filesystem::path p = "example.txt";
6
7     // Check if the file exists
8     if (std::filesystem::exists(p)) {
9         std::cout << p << " exists." << std::endl;
10    } else {
11        std::cout << p << " does not exist." << std::endl;
12    }
13
14    // Create a new directory
15    std::filesystem::create_directory("new_directory");
16
17    // Remove a file
18    std::filesystem::remove("example.txt");
19
20    // Rename a file
21    std::filesystem::rename("old_name.txt", "new_name.txt");
22
23    // Copy a file
24    std::filesystem::copy("source.txt", "destination.txt");
25
26    return 0;
27 }
```

Iterating Over Directory Contents

You can use `std::filesystem::directory_iterator` to iterate over the contents of a directory.

② Example:

```
1 #include <iostream>
2 #include <filesystem>
3
4 int main() {
5     std::filesystem::path dir = "some_directory";
6
7     // Iterate over the contents of the directory
8     for (const auto& entry : std::filesystem::directory_iterator(
9         dir)) {
10        std::cout << entry.path() << std::endl;
11    }
12
13    return 0;
14 }
```

Recursive Directory Iteration

You can use `std::filesystem::recursive_directory_iterator` to iterate over the contents of a directory recursively.

?

Example:

```
1 #include <iostream>
2 #include <filesystem>
3
4 int main() {
5     std::filesystem::path dir = "some_directory";
6
7     // Recursively iterate over the contents of the directory
8     for (const auto& entry : std::filesystem::
9         recursive_directory_iterator(dir)) {
10         std::cout << entry.path() << std::endl;
11     }
12
13 }
```

Download

5 CMake

5.1 Introduction

CMake stands for "Cross-Platform Make." It is a **build-system generator**, meaning it creates the files (e.g., `Makefile`, Visual Studio project files) needed by your build system to compile and link your project. CMake abstracts away platform-specific build configurations, making it easier to maintain code that needs to run on multiple platforms.

It works the following way:

1. You write a `CMakeLists.txt` file that describes your project's configuration and structure.
2. You run CMake on the `CMakeLists.txt` file to generate the build system files (e.g. `Makefile` on Linux or `.sln` for Visual Studio).
3. You use the generated build system to compile and link your project.

5.2 CMakeLists.txt

Contains the configuration and structure of your project. It is a script that CMake uses to generate the build system files. It has the following structure:

```
cmake_minimum_required(VERSION 3.10)
project(MyProject)

add_executable(my_project_main.cpp)
```

5.2.1 Minimum Version

Here is the first line of every `CMakeLists.txt`, which is the required name of the file CMake looks for:

```
cmake_minimum_required(VERSION 3.10)
```

The version on CMake dictates the policies. Starting in CMake 3.12, this supports a range like `3.12 ... 3.15`. This is useful when you want to use new features but still support older versions.

```
cmake_minimum_required(VERSION 3.12...3.15)
```

5.2.2 Setting a project

Every top-level CMake file will have this line:

```
project(MyProject VERSION 1.0
        DESCRIPTION "My Project"
        LANGUAGES CXX)
```

Strings are quoted, whitespace does not matter and the name of the project is the first argument. All the keywords are optional. The `version` sets a bunch of variables, like `MyProject_VERSION` and `PROJECT_VERSION`. The `LANGUAGES` keyword sets the languages that the project will use. This is useful for IDEs that support multiple languages.

5.2.3 Making an executable

```
add_executable(my_project_main my_project_main.cpp)
```

`my_project` is both the name of the executable file generated and the name of the CMake target created. The source file comes next and you can add more than one source file. CMake will only compile source file extensions. The headers will be ignored for most purposes; they are there only to be shown up in IDEs.

5.2.4 Making a library

```
add_library(my_library STATIC my_library.cpp)
```

`STATIC` is the type of library. It can be `SHARED` or `MODULE`. The source files are the same as for executables. Often you'll need to make a fictional target, i.e., one where nothing needs to be compiled, for example for header-only libraries. This is called an `INTERFACE library`, and the only difference is that it cannot be followed by filenames.

5.2.5 Targets

Now we've specified a target, we can set properties on it. CMake is all about targets and properties. An executable is a target, a library is a target. Your application is built as a collection of targets depending on each other.

```
target_include_directories(my_library PUBLIC include)
```

This sets the include directories for the target. The `PUBLIC` keyword means that the include directories will be propagated to any target that links to `my_library`. We can then chain targets:

```
add_library(my_library STATIC my_library.cpp)
target_link_libraries(my_project PUBLIC my_library)
```

This will link `my_project` to `my_library`. The `PUBLIC` keyword means that the link will be propagated to any target that links to `my_project`.

Targets can have include directories, linked libraries (or linked targets), compile options, compile definitions, compile features and more.

5.2.6 Variables

`Local variables` are used to store values that are used only in the current scope:

```
set(MY_VAR "some_file")
```

The names of the variables are case-sensitive and the values are strings. You access a variable by using `$${}`. CMake has the concept of scope; you can access the value of the variable after you set it as long as you are in the same scope. If you leave a function or a file in a sub directory, the variable will no longer be defined. You can set a variable in the scope immediately above your current one with `PARENT_SCOPE` at the end.

One can also set a list of values:

```
set(MY_LIST "value1" "value2" "value3")
```

which internally becomes a string with semicolons. You can access the values with `${MY_LIST}` . If you want to set a variable from the command line, CMake offers a variable cache. `Cache variables` are used to interact with the command line:

```
set(MY_CACHE_VAR "VALUE" CACHE STRING "Description")
option(MY_OPTION "Set from command line" ON)
```

Then:

```
cmake /path/to/src/ \
-DMY_CACHE_VAR="some_value" \
-DMY_OPTION=OFF
```

`Environment variables` are used to interact with the environment:

```
# Read
message(STATUS $ENV{MY_ENV_VAR})

# Write
set(ENV{MY_ENV_VAR} "some_value")
```

But it is not recommended to use environment variables in CMake.

5.2.7 Properties

The other way to set properties is to use the `set_property` command:

```
set_property(TARGET my_library PROPERTY CXX_STANDARD 17)
```

This is like a variable, but it is attached to a target. The `PROPERTY` keyword is optional. The `CXX_STANDARD` is a property that sets the C++ standard for the target.