



UniTs - University of Trieste

Faculty of Scientific and Data Intensive Computing
Department of mathematics informatics and geosciences

Advanced Programming for Astrophysics

Lecturers:

Prof. Murante Giuseppe
Prof. Taffoni Giuliano

Author:

Andrea Spinelli

October 21, 2025

This document is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike \(CC BY-NC-SA\)](#) license. You may share and adapt this material, provided you give appropriate credit, do not use it for commercial purposes, and distribute your contributions under the same license.

Preface

This document was prepared for the Advanced Programming for Astrophysics course, held by Prof. Murante Giuseppe and Prof. Taffoni Giuliano in the academic year 2025/2026.

Contents

1	Introduction	1
1.1	The Unix Shell	2
1.1.1	Basic Commands	2
1.1.2	Environment Variables	4
1.1.3	File Permissions	4
2	Intro to C	6
2.1	Compilation	6
2.1.1	Libraries	7
2.1.2	Data Types	7
2.1.3	Directives	7
2.1.4	Functions	8
2.1.5	Variable Types	9
2.1.6	Operators	10
2.2	Control Structures	10
2.2.1	Conditional statements	11
2.2.2	Loops	11
2.3	Memory Management	14
2.3.1	Memory Layout	14
2.3.2	Pointers	15
2.3.3	Command-Line Arguments	18
2.3.4	Store Management	19
2.3.5	Linked Lists	20
2.3.6	Binary Trees	21
2.3.7	Hash Tables	22

1

Introduction

Operating systems (OS) are the core software that manage a computer's hardware and software resources. Among the most widely used are Microsoft Windows, Apple's macOS, and the open-source Linux. While they differ in many aspects, both macOS and Linux are part of the ***Unix-like*** family of operating systems. These systems are renowned for their stability, security, and a powerful command-line interface known as the ***shell***.

This chapter provides an introduction to the Unix shell, exploring its fundamental role in interacting with the system.

File System

A crucial component of the operating system is the file system, which organizes files and directories on a computer. It is a crucial component of the operating system that allows users to store, retrieve, and manage data efficiently. In unix-like systems, the file system is organized in a hierarchical structure.




Figure 1.1: The File System

Path

A path specifies the unique location of a file or directory within the file system's hierarchy. Paths can be ***absolute***, starting from the root directory (`/`), or ***relative*** to the current working directory.

<code>/</code>	Root Directory
<code>.</code>	Current Directory
<code>~</code>	Home Directory
<code>..</code>	Parent Directory

?

Example: *Example of a path*

- a file in the user's Desktop: `/home/user/Desktop/file.txt` or `~/Desktop/file.txt`
- a file in the current directory: `./file.txt`
- a file in the parent directory: `../file.txt`
- a file in the home directory: `~/file.txt`

1.1 The Unix Shell

What is a shell?

The shell is the primary interface between users and the computer's core system. When you type commands in a terminal, the shell interprets these instructions and communicates with the operating system to execute them, making complex system operations accessible through simple commands.

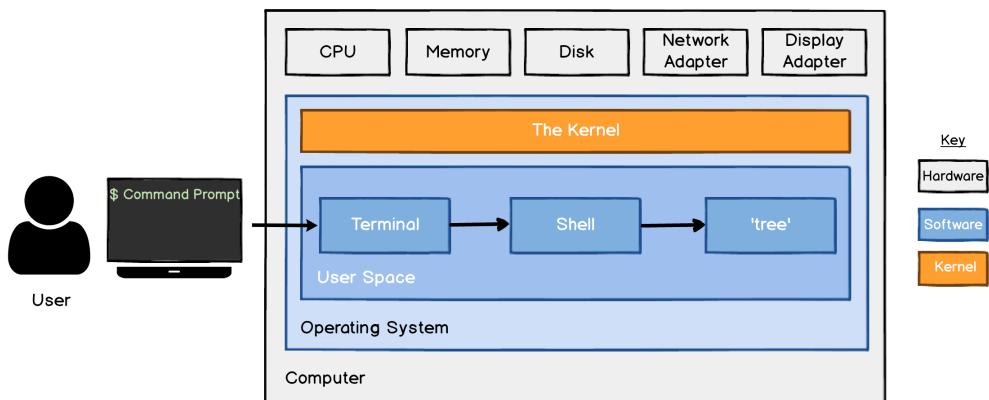


Figure 1.2: The Unix Shell

Definition: *Shell*

A shell is a program that provides the traditional, text-only user interface for Linux and other UNIX-like operating systems. Its primary function is to read commands that are typed into a console [...] and then execute (i.e., run) them. The term shell derives its name from the fact that it is an outer layer of an operating system. A shell is an interface between the user and the internal parts of the OS (at the very core of which is the kernel).

~Linfo

1.1.1 Basic Commands

A command is nothing more than a program that is executed by the shell. The usual syntax is:

```
command [options] [arguments]
```

The options can be passed setting the corresponding flag e.g. `-h` for help.

Navigation Commands

It is possible to navigate and manage files or directories using the following commands:

Command	Description
<code>pwd</code>	Print Working Directory
<code>cd</code>	Change Directory
<code>ls</code>	List Directory
<code>mkdir</code>	Make Directory
<code>rm</code>	Remove File or Directory
<code>cp</code>	Copy File or Directory
<code>mv</code>	Move File or Directory
<code>touch</code>	Create File

Recursive Commands

`rm`, `cp` and many other commands can be used recursively into a folder using the flag `-r` (recursive).

Help Flag

Many commands support the `--help` flag to display information about the command and its usage.

💡 Tip: Hidden Files

Hidden files are files that are not shown by default when listing the contents of a directory. They are usually marked with a dot (.) at the beginning of their name. e.g. `.config`

Searching Tools

Unix-like systems provide powerful tools for searching files and content. These tools often support regular expressions (regex), which are patterns used to match character combinations in text.

File Search Commands

Command	Description	Example
<code>find</code>	Search for files and directories	<code>find /home -name "*txt"</code>
<code>locate</code>	Fast file search using database	<code>locate filename</code>
<code>grep</code>	Search text within files	<code>grep "pattern" file.txt</code>
<code>which</code>	Find location of executable	<code>which gcc</code>
<code>whereis</code>	Find binary, source, and manual pages	<code>whereis python</code>

Wildcards and Pattern Matching

Wildcards are special characters that represent one or more characters in a filename or text string:

Character	Name	Description
*	asterisk	Matches zero or more characters
?	question mark	Matches exactly one character
[]	brackets	Matches any single character within brackets
[!]	negated brackets	Matches any character not in brackets

💡 Example: Wildcard Examples

- `*.txt` : matches all files ending with .txt
- `file?.c` : matches file1.c, file2.c, but not file10.c
- `[abc]*` : matches files starting with a, b, or c
- `[!0-9]*` : matches files not starting with digits

Operators

Commands can be combined using operators to manipulate input and output streams:

Operator	Name	Description
>	output redirection	Redirects standard output to a file, overwriting its contents
>>	append redirection	Redirects standard output, appending it to the end of a file
	pipe	Passes the output of a command as input to another
&&	logical AND	Executes the next command only if the previous one succeeds
	logical OR	Executes the next command only if the previous one fails
;	separator	Executes commands sequentially, regardless of their outcome

1.1.2 Environment Variables

Environment variables are dynamic values that affect the behavior of processes and programs running on the system. They provide a way to pass configuration information to applications without hardcoding values into the program. These variables are stored in the system's environment and can be accessed by any process.

Variable	Description	Example
PATH	Colon-separated list of dirs to search for executables	/usr/bin:/bin:/usr/sbin
HOME	Path to the user's home directory	/home/username
USER	Current username	username
SHELL	Path to the current shell	/bin/bash
PWD	Current working directory	/home/username/Desktop

To view environment variables, use the `env` command or `echo $VARIABLE_NAME`. To set a variable temporarily, use `export VARIABLE_NAME=value`. For permanent changes, modify configuration files like `.bashrc` or `.bash_profile`.

1.1.3 File Permissions

Unix-like systems use a permission system to control access to files and directories. Each file has three types of permissions: read (r), write (w), and execute (x). These permissions are assigned to three categories of users: owner (user), group, and others.

Permission Structure

File permissions are displayed using a 10-character string where the first character indicates the file type, and the remaining nine characters represent permissions for owner, group, and others:

Character	Position	Meaning
-	1st	Regular file
d	1st	Directory
r	2nd, 5th, 8th	Read permission
w	3rd, 6th, 9th	Write permission
x	4th, 7th, 10th	Execute permission

For instance, `-rwxr-xr--` means:

- Regular file (-)
- Owner: read, write, execute (rwx)
- Group: read, execute (r-x)
- Others: read only (r-)

Octal Notation

Permissions can be represented numerically using octal notation:

(r) Read	4	100
(w) Write	2	010
(x) Execute	1	001

By summing these values for each user category (owner, group, and others), we get a three-digit code. For example, `rwx` permissions translate to $4 + 2 + 1 = 7$, and `r-x` to $4 + 0 + 1 = 5$. Thus, a full permission string like `rwxr-xr-x` can be concisely represented as 755.

Common combinations include: **755** (`rwxr-xr-x`), **644** (`rw-r--r--`), and **600** (`rw----`).

Permission Commands

Command	Description	Example
<code>chmod</code>	Change file permissions	<code>chmod 755 script.sh</code>
<code>chown</code>	Change file ownership	<code>chown user:group file.txt</code>
<code>chgrp</code>	Change group ownership	<code>chgrp staff file.txt</code>
<code>umask</code>	Set default permissions	<code>umask 022</code>

TODO: notes about file viewing and text processing, few words about multiprocessing and context switching

2

Intro to C

C is a general-purpose programming language created in the 1970s by Dennis Ritchie. Its design gives programmers direct access to the underlying hardware, making it highly efficient and powerful. It is instrumental in implementing operating systems, device drivers, and protocol stacks.

The easiest code we can write in C is:

```
1 #include <stdio.h>
2
3 int main() {
4     printf("Hello, World!\n");
5 }
```

- `#include <stdio.h>`: is a preprocessor directive.
- `int main()`: is the main function.
 - `int`: is the return type of the function, if omitted, the function returns an integer.
 - `main()`: is the name of the function.
- `printf(" ... ")`: is the function that prints the string `"Hello, World!"` to the console.
- `{}`: is the block of code that contains the statements of the function.

2.1 Compilation

To compile a C program, use the following command:

```
gcc main.c -o main.x
```

The compiler generates the object code, which is a binary file that contains the machine language translation of the program. The command above actually executes 3 steps:

1. **Preprocessing**: This step expands the macros and includes the header files. e.g.
`#include <stdio.h>`

```
gcc -E main.c -o main_preprocessed.c
```

2. **Compilation**: The compiler translates the preprocessed code into object code.

```
gcc -c main_preprocessed.c -o main.o
```

During this step, it is possible to set the `-I` flag to include other directories (useful if the headers are not in the default directory). e.g. `-I /path/to/include`

3. **Linking**: The linker combines the object code with the standard library to create an executable.

```
gcc main.o -o main.x
```

The `-o` flag is used to set the output file name.

During this step, it is possible to set the `-L` flag to include other directories (useful if the libraries are not in the default directory). e.g. `-L /path/to/lib` and to link against external libraries using the `-l` flag. e.g. `-lmylib`

2.1.1 Libraries

A library is a collection of pre-compiled code, such as functions and data structures, that can be reused across multiple programs. They promote code modularity and prevent developers from having to reimplement common functionalities, such as input/output operations or mathematical calculations. There are two types of libraries:

- `lib__.a` : is a static library.
- `lib__.so` : is a shared library.

The difference between the two is that the static library is linked directly into the executable at compile time, while the shared library is loaded at runtime.

2.1.2 Data Types

C provides several fundamental data types to represent different kinds of data. Understanding these types is crucial for writing efficient and correct programs.

TODO: notes about data types

Increment and Decrement

The `++` and `--` operators are used to increment and decrement the value of a variable.

There are two types of increment and decrement operators:

- `+i` (`--i`): is the prefix increment (decrement) operator. The value of the variable is incremented (decremented) before it is used.
- `i++` (`i--`): is the postfix increment (decrement) operator. The value of the variable is incremented (decremented) after it is used.

2.1.3 Directives

Directives are used to control the compilation process. They are used to define constants, macros, and to conditionally compile code.

Macros

Macros are used to define values that are not expected to change during the execution of the program. They are defined using the `#define` directive.

```
1 #define CYCLE 0
```

Conditional Compilation

Conditional compilation is used to compile code only if a certain condition is met. They are defined using the `#ifdef`, `#ifndef`, `#else`, `#elif`, and `#endif` directives.

```
#ifdef CYCLE
    // code to compile if CYCLE is defined
#endif
```

2.1.4 Functions

Functions are used to group code into *routines*: reusable units of code.

```
1 return-type function-name(parameters declarations, if any) {  
2     declarations  
3     statements  
4     return return-value;  
5 }
```

Parameters are passed to the function by value, meaning that a copy of the argument is passed to the function. If the argument is a *pointer* (we will see), the function will modify the original argument.

It is possible to return only a single value. We will see later that it is possible to return a *pointer* or a *struct*, which can be useful to return multiple values from a function.

If the return value is not the same specified in the function declaration, the compiler probably won't notice, so be careful.

Variable scopes

- **Local scope:**

Variables declared inside a function are said to have *local scope*. This means they only exist and are accessible from within that specific function. Once the function finishes its execution, these variables are destroyed. This allows different functions to use the same variable names without causing conflicts.

- **Global scope:**

Variables declared outside any function are said to have *global scope*. This means they are accessible from any function in the program.

⌚ Observation: Variable shadowing

If a local variable is declared with the same name as a global variable, the local variable takes precedence within its scope. This is known as *variable shadowing*. The global variable is temporarily hidden, and any reference to that variable name inside the function will refer to the local one. The global variable remains unaffected outside of this scope.

```
1 #include <stdio.h>  
2  
3 int a = 10; // Global variable  
4  
5 void myFunction() {  
6     int a = 20; // Local variable shadows the global one  
7     printf("Inside function, a = %d\n", a); // Prints 20  
8 }  
9  
10 int main() {  
11     printf("Before function call, a = %d\n", a); // Prints 10  
12     myFunction();  
13     printf("After function call, a = %d\n", a); // Prints 10  
14     return 0;  
15 }
```

2.1.5 Variable Types

In C, every variable has a type, which dictates the size and layout of its memory, the range of values it can store, and the set of operations that can be applied to it. The fundamental types are used to represent integers, floating-point numbers, and characters. The table below shows common data types and their typical sizes on modern systems.

Type	Typical Size	Description
<code>char</code>	1 byte	Stores a single character or a small integer.
<code>short int</code>	2 bytes	Short integer.
<code>int</code>	4 bytes	The most natural size of integer for the machine.
<code>long int</code>	8 bytes	Long integer.
<code>float</code>	4 bytes	Single-precision floating-point number.
<code>double</code>	8 bytes	Double-precision floating-point number.

Tip: *implementation-defined sizes*

Note that the exact sizes of these types are implementation-defined and can vary across different systems and compilers. You can use the `sizeof` operator to determine the size of a type on your machine.

Integer types (`char`, `short`, `int`, `long`) can be either `signed` (the default) or `unsigned`. A `signed` type can hold both positive and negative values, while an `unsigned` type can only hold non-negative ones. By using the sign bit to store value instead, an `unsigned` type can represent a maximum value twice as large as its signed counterpart. For example, a `signed char` typically ranges from -128 to 127, whereas an `unsigned char` ranges from 0 to 255.

Type conversion

C automatically converts the smaller type to the larger type.

```
int x = 10;
float y = 20.5;
float z = x + y; // z = 30.5
```

For the other cases, we need to use an explicit cast.

```
1 double x = 10.;
2 int y = 7;
3 int z = y + (int) x; // z = 17
```

Declaration and initialization

```
1 // only declaration
2 int x, y, z;
3 char c, line[1000];
4
5 // declaration and initialization
6 int x = 10, y = 20, z = 30;
7 char line[1000] = "Hello"; // or = { 'H', 'e', 'l', 'l', 'o', '\0' };
```

2.1.6 Operators

C provides a rich set of operators, which are summarized in the table below.

Category	Operator	Description
Arithmetic	<code>+</code>	Addition
	<code>-</code>	Subtraction
	<code>*</code>	Multiplication
	<code>/</code>	Division
	<code>%</code>	Modulo (remainder of division)
Relational	<code>==</code>	Equal to
	<code>!=</code>	Not equal to
	<code>></code>	Greater than
	<code><</code>	Less than
	<code>≥</code>	Greater than or equal to
	<code>≤</code>	Less than or equal to
Logical	<code>&&</code>	Logical AND
	<code> </code>	Logical OR
	<code>!</code>	Logical NOT
Bitwise	<code>&</code>	Boolean AND
	<code> </code>	Boolean OR
	<code>^</code>	Boolean XOR
	<code>~</code>	Boolean NOT
	<code><<</code>	Left shift (\equiv int mul. by 2)
	<code>>></code>	Right shift (\equiv int div. by 2)

It is possible to use the compound assignment operators (`+=`, `-=`, ...), which are equivalent to the corresponding arithmetic operator followed by the assignment operator.

```
int x = 10;  
x += 5; // x = 15
```

⚠ Warning: Precedence

Operations have a precedence, which is used to determine the order in which they are evaluated. It is possible to change the order of evaluation using parentheses.

2.2 Control Structures

Control structures are used to control the flow of the program. They are used to execute a block of code only if a certain condition is met.

2.2.1 Conditional statements

If-Else

```
1 if (condition1) {  
2     // code  
3 }  
4 else if (condition2) {  
5     // code  
6 }  
7 else { // any other condition  
8     // code  
9 }
```

It is possible to nest if statements one inside the other. If you do, make sure to use braces to avoid ambiguity.

Switch

The `switch` statement is used to execute different blocks of code based on the value of a variable.

```
1 switch (expression) {  
2     case constant1:  
3         // code  
4     case constant2:  
5         // code  
6     default: // any other value  
7         // code  
8 }
```

The default behaviour of the `switch` statement is to check all the cases, even if the expression matches one of them. To avoid this, it is possible to use the `break` statement to exit the switch statement.

```
1 switch (expression) {  
2     case constant1:  
3         // code  
4         break;  
5     case constant2:  
6         // code  
7         break;  
8     default:  
9         // code  
10        break; // implicit  
11 }
```

2.2.2 Loops

Loops are used to execute a block of code a specified number of times. There are two main types of loops: the `while` loop and the `for` loop.

While and Do While Loops

The `while` loop is used to execute a block of code as long as a specified condition is true.

```
1 while (condition) {  
2     // code  
3 }
```

If the condition is never met, the block of code is not executed. If we need to execute the block of code at least once, we can use the `do while` loop.

```
1 do {  
2     // code  
3 } while (condition);
```

For Loop

The `for` loop is used to execute a block of code a specified number of times.

```
1 for (int i = 0; i < 5; i++) {  
2     // code  
3 }
```

Break and Continue

The `break` statement is used to exit a loop prematurely. The `continue` statement is used to skip the rest of the code in a loop and move to the next iteration.

```
1 while (1) { // infinite loop (should be avoided)  
2     // code  
3     if (condition_1) {  
4         break; // exit the loop  
5     }  
6     else if (condition_2) {  
7         continue; // skip the code below and move to the next iteration  
8     }  
9 }
```

MISSING: 03/10/25 notes

10/10

2.3 Memory Management

The C language allows programmers to manage memory manually, providing powerful capabilities for fine-grained control over how memory is allocated, used, and freed. This flexibility enables efficient use of system resources, but also requires careful attention to avoid errors such as memory leaks or accessing invalid memory.

2.3.1 Memory Layout

The memory layout of a program (Figure 2.1) is the way the memory is organized in the computer.

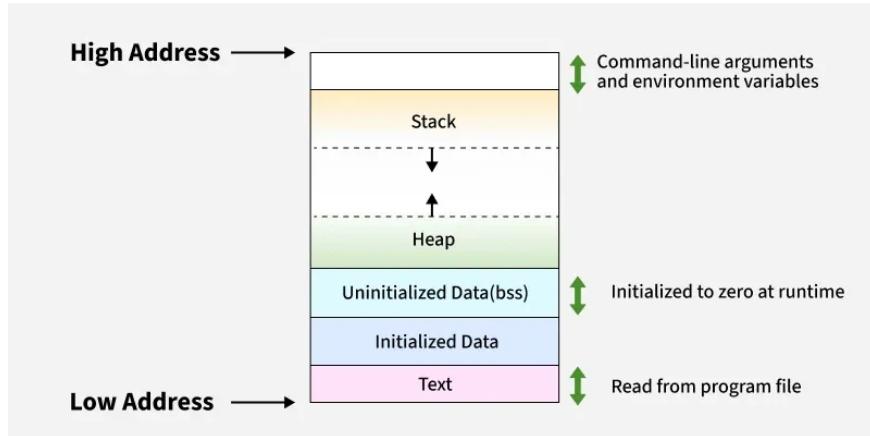


Figure 2.1: The Memory Layout

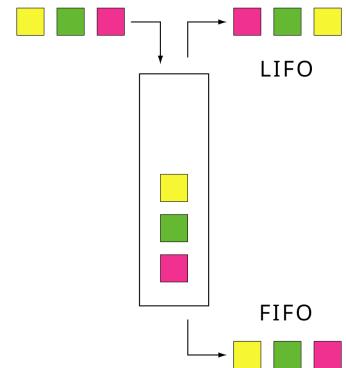


Figure 2.2:
FIFO and LIFO

- **Stack**: It is used to store *local variables*, *function parameters*, and *return addresses*. Whenever a function is called, a new *stack frame* is created to hold its local data. The stack automatically grows and shrinks as functions are called and return. Variables stored here are automatically removed when the function finishes, following a Last In, First Out (LIFO) order.
- **Heap**: It is used for dynamic memory allocation, which means memory that is explicitly requested at runtime (using functions like `malloc` and `free` in C). Memory allocated on the heap persists until it is explicitly freed by the programmer, and is not automatically cleaned up.
- **BSS (Block Started by Symbol)**: It contains all global and static variables that are declared but not initialized by the programmer. The operating system initializes this region to zero before the program starts running.
- **Data**: It stores global and static variables that are explicitly initialized by the programmer. For example, `int x = 5;` at global scope would be stored here. The values in this segment are set to their initial values when the program starts.
- **Text**: Also known as the code segment, it contains the compiled machine code instructions of the program itself. It is typically marked as read-only to prevent accidental modification of the program's instructions during execution.

⚠ Warning: Reserved addresses

Low addresses are reserved for the kernel and other critical system components. Do not use them for your own variables if you don't know what you are doing.

2.3.2 Pointers

All data, variables, and functions in a C program reside in the computer's memory. Each location in memory is identified by a unique address, typically represented as a hexadecimal value. These addresses allow the program to access and manipulate data efficiently. Understanding how to work with memory addresses is fundamental in C, as it enables direct interaction with the underlying hardware and forms the basis for concepts like pointers, dynamic memory allocation, and efficient data structures.

Consider the following code and its corresponding memory diagram:

```
1 int a = 1;           // a is 1
2 a++;                // a is now 2
3
4 int* p;             // p is a pointer to an integer
5 p = &a;              // p now points to a
6
7 int value = *p; // value is now 2
```

Memory	Address
a = 2	0x0
p = 0x0	0x1
value = 2	0x2
:	:

Explanation:

- `a` is an integer variable stored at address `0x0`, and after `a++`, its value is 2.
- `p` is a pointer variable stored at address `0x1`, and it holds the address of `a` (`0x0`).
- `value` is an integer variable stored at address `0x2`, and it is assigned the value pointed to by `p`, which is 2.

The `&` operator is used to get the address of a variable (e.g., `p = &a;`). The `*` operator is used to access the value stored at the address pointed to by a pointer (e.g., `value = *p;`).

Pointers arithmetic

We have seen `*` and `&` operators, used to access the value stored at the address pointed to by a pointer and to get the address of a variable, respectively. Now, let's see how pointers behaves with arithmetic operations.

The actual size of the increment (decrement) is the size of the type pointed to by the pointer.

```
1 int a = 1;
2 int *p = &a;
3 p++; // this is actually p + 4 bytes
4
5 double b = 2.0;
6 double *q = &b;
7 q++; // this is actually q + 8 bytes
```

The same applies to the decrement, addition, subtraction, multiplication and division operators (the last two are rarely used with pointers).

We have seen that to access the value pointed to by the pointer, we can use the `*` operator. It is possible to perform arithmetic operations among the values pointed by pointers using the same operator.

```
1 int a = 1;
2 int *p = &a;
3 (*p) += 2; // a is now 3
```

Pointers and Function Arguments

When a function is called, the arguments are passed by value, meaning that a copy of the argument is passed to the function. If the argument is a pointer, the function can modify the original argument:

```
1 void swap(int *a, int *b) {  
2     int temp = *a;  
3     *a = *b;  
4     *b = temp;  
5 }
```

The code above swaps the values of the two variables.

Pointers and Arrays

In C, arrays and pointers are closely related. An array name is essentially a pointer to the first element.

```
1 int a[5]; // only declaration  
2 int a[5] = {1, 2, 3, 4, 5}; // declaration and initialization  
3 int *ptr = a; // pointer to the first element
```

During the declaration, the number in the square brackets is the size of the array.

The expressions `a[i]` and `*(a + i)` are equivalent. Let's see how to

```
1  
2 for (int i = 0; i < 5; i++) {  
3     a[i]++;
4 }
5  
6 for (int i = 0; *(a + i) != 0; i++) {
7     (*a + i)++;
8 }
9  
10 * missing one *
```

The two loops perform the same operation. (The third one, missing, do the same but in a more efficient way - slightly noticeable for "short arrays")

Observation: Strings and Pointers

Strings and pointers are basically the same thing: a string is an array of characters, ending with the null character `'\0'`.

The code below copies the string `t` into the string `s` in a compact and elegant way.

```
1 void strcpy(char *s, char *t) {  
2     while (*s++ = *t++);  
3 }
```

Multidimensional Arrays

Multidimensional arrays are arrays of arrays. They are declared using multiple square brackets.

```
1 int a[3][4] = {  
2     {1, 2, 3, 4},  
3     {5, 6, 7, 8},  
4     {9, 10, 11, 12}  
5 };
```

The code above declares a 2D array of integers with 3 rows and 4 columns.

The expressions `a[i][j]` and `*(&a[i] + j)` are equivalent.

💡 Tip: Declaring pointers and 2D arrays

There are several ways to declare variables related to 2D arrays in C, each with different meanings and use cases:

```
int a[3][4];    // a true 2D array: 3 rows, 4 columns  
int *p[4];      // an array of 4 pointers to int  
int **q;        // a pointer to a pointer to int
```

Explanation:

- `int a[3][4];` declares a contiguous block of memory for 12 integers (3 rows, 4 columns). This is the standard way to declare a 2D array.
- `int *p[4];` declares an array of 4 pointers to `int`. Each element of `p` can point to the beginning of a separate row (or any `int`), but the rows themselves are not stored contiguously unless you arrange them that way.
- `int **q;` declares a pointer to a pointer to `int`. This is often used for dynamically allocated 2D arrays, where both the rows and the columns can be allocated at runtime.

Key differences:

- The first declaration (`a`) allocates all the memory for the array at once and ensures the data is stored contiguously.
- The second (`p`) and third (`q`) declarations only allocate space for pointers, not for the actual integers. You must allocate memory for the data separately, typically using `malloc`.
- Only the first declaration (`a`) knows the size of both dimensions at compile time.
We will explore dynamic memory allocation for 2D arrays in more detail later.

2.3.3 Command-Line Arguments

In C, the `main` function can be defined to accept arguments from the command line, allowing users to pass information to your program when it starts.

```
1 int main(int argc, char *argv[]) { // Your code here }
```

- `argc` (argument count) is an integer representing the number of command-line arguments passed to the program, including the program's name itself.
- `argv` (argument vector) is an array of pointers to strings (character arrays), where each string is one of the command-line arguments.

Important details:

- `argv[0]` is always the name (or path) of the program as it was invoked.
- The actual user-supplied arguments start from `argv[1]` up to `argv[argc-1]`.
- `argc` is always at least 1, since the program name is always present.

?

Example: Example

If you run the program as follows:

```
$ ./myprog 1 -2 1
```

Then:

- `argv[0]` is `./myprog`
- `argv[1]` is `"1"`
- `argv[2]` is `"-2"`
- `argv[3]` is `"1"`

⚠ Warning: Arguments as strings

Arguments passed to the program are treated as strings. You need to convert them to the appropriate type using the `atoi`, `atof`, `atol`, `atoll` functions.

You can use a loop to process all arguments:

```
1 for (int i = 0; i < argc; i++) {
2     printf("Argument %d: %s\n", i, argv[i]);
3 }
```

This feature is especially useful for writing flexible and user-friendly command-line programs.

Pointers to functions

Pointers to functions are used to store the address of a function. They are declared using the `typedef` keyword.

```
typedef int (*func_ptr)(int, int);
```

The code above declares a pointer to a function that takes two `int` arguments and returns an `int`.

MISSING: Variable lenght argument list

2.3.4 Store Management

It is possible to allocate (and deallocate) memory dynamically. To do that we use functions such as `malloc`, `calloc` and `free`:

- `void *malloc(size_t size)`: returns a pointer to `size` bytes of uninitialized storage, or `NULL` if the request cannot be satisfied.
- `void *calloc(size_t n, size_t size)`: returns a pointer to *contiguous* storage for `n` elements of `size` bytes each, or `NULL` if the request cannot be satisfied. *The memory is initialized to zero.*

The pointer returned by these functions is a `void *` and we need to cast it to the appropriate type. When a dynamically allocated pointer is no longer used, release the memory with `free`.

```
1 /* allocate n integers */
2 size_t n = 100;
3 int *values = malloc(n * sizeof *values);
4 if (values == NULL) {
5     /* handle error */
6 }
7
8 /* ... use values[0..n-1] ... */
9
10 /* alternatively, get zero-initialized memory */
11 int *zeros = calloc(n, sizeof *zeros);
12 if (zeros == NULL) {
13     free(values);
14     /* handle error */
15 }
16
17 /* ... use zeros and values ... */
18
19 free(zeros);
20 free(values);
```

💡 Tip: Free order

It's good practice to free dynamically allocated memory in the reverse order that you allocated it. This helps avoid bugs, especially if your allocations depend on each other. While many modern compilers can handle this automatically, freeing memory in the wrong order can sometimes lead to problems like segmentation faults.

2.3.5 Linked Lists

A **linked list** is a *dynamic* data structure that consists of a sequence of nodes, where each node contains a *value* and a *pointer* to the next node.

There are also **doubly linked lists**, where each node contains a value and pointers to both the next and the previous node, allowing traversal in both directions.

```
typedef struct Node {
    int data;
    struct Node *next;
    // doubly linked list
    // struct Node *prev;
};
```

Here are some functions to create, print and free a linked list (to generalize, I made some of them iterative, some recursive):

```
1 // Return a pointer to a new node with the given data
2 Node* create_node(int data) {
3     Node *newNode = (Node*) malloc(sizeof(Node));
4     if (newNode == NULL) {
5         printf("Error: unable to allocate memory.\n");
6         exit(1);
7     }
8     newNode->data = data;
9     newNode->next = NULL;
10    return newNode;
11 }
12
13 // Insert a new node at the beginning of the list
14 void insert_at_head(Node **head, Node *node_to_insert) {
15     node_to_insert->next = *head;
16     *head = node_to_insert;
17 }
18
19 // Print the list
20 void print_list(Node *head) {
21     for (Node *curr = head; curr != NULL; curr = curr->next)
22         printf("%d ", curr->data);
23     printf("\n");
24 }
25
26 // Free the list
27 void free_list(Node *head) {
28     if (head == NULL) return;
29     free_list(head->next);
30     free(head);
31 }
```

⚠ Warning: Memory management

Always `free` the nodes you allocate when you no longer need them to avoid memory leaks.

Advantages and disadvantages

Pros: The size of a linked list can be changed dynamically at runtime. Inserting and deleting elements is efficient, especially in doubly linked lists.

Cons: Access to the n -th element is slow, as you need to traverse the list from the beginning. Each node also requires extra memory for its pointer(s), which increases memory usage.

2.3.6 Binary Trees

A **binary tree** is a hierarchical data structure in which each node contains a *data* and has at most two child nodes, typically referred to as the *left* and *right* child. Each node also stores pointers to these children, creating a branching structure that is fundamentally different from the linear nature of arrays or linked lists.

A particularly common type is the **Binary Search Tree** (BST), which imposes an ordering: for any given node, all values in its left subtree are less, and all values in its right one are greater. This property allows for efficient searching, insertion, and deletion operations, typically in $O(\log n)$ for balanced trees.

A binary tree in C can be defined with a structure where each node stores data and pointers to its left and right children.

The entire tree is represented by a pointer to its root node. If it is empty, its pointer is set to `NULL`.

To process all the nodes in a binary tree, **recursive** traversal is widely used. The main orders are:

- **Inorder**: Visit the left subtree, then the root, then the right subtree.
- **Preorder**: Visit the root, then the left subtree, then the right subtree.
- **Postorder**: Visit the left subtree, then the right subtree, then the root.

Below is an example of how to add a new element to a binary search tree:

```
1 struct Node* add(struct Node* p, int v) {
2     if (p == NULL) { // If the tree is empty, create a new node
3         p = (struct Node *) malloc(sizeof *p);
4         p->data = v;
5         p->left = p->right = NULL;
6     }
7     // If the value is less than the current node, go left
8     else if (v < p->data) {
9         p->left = add(p->left, v);
10    }
11    // If the value is greater than the current node, go right
12    else if (v > p->data) {
13        p->right = add(p->right, v);
14    } else { /* The value is already present, do nothing */ }
15    return p;
16 }
```




Figure 2.3: Binary search tree

```
struct Node {
    int data;
    struct Node *left;
    struct Node *right;
};
```

Advantages and disadvantages

Pros: Binary search trees (especially if balanced) allow for fast operations: searching, insertion, and deletion are all $O(\log n)$ on average. They also keep data automatically sorted, which makes ordered traversals straightforward.

Cons: If the tree becomes unbalanced (e.g., when many values are inserted in sorted order), performance degrades to $O(n)$. Extra care or specific algorithms are needed to keep it balanced.

2.3.7 Hash Tables

Suppose we have a collection of elements, each associated with a unique key, and we want to efficiently store and retrieve them by key—ideally in constant time. A powerful solution for this is the **hash table**.

A **hash table** is a data structure that organizes data based on a computed index, so that insertion, deletion, and lookup operations can be performed very efficiently (typically with average-case $O(1)$ time complexity). Instead of searching through a list, we use a *hash function* to compute an index (or "address") in an array (the "table") where the element should be stored.

The basic idea is as follows:

- A **hash function** transforms a key (which can be a string, number, etc.) into an integer, which is then mapped to one of the available array indices.
- Each position in the array is called a **bucket** or **slot**.
- Ideally, each key maps to a unique index. However, since there are more possible keys than buckets, multiple keys may map to the same index, a situation known as a **collision**.
- To handle collisions, each bucket can store a list (or another structure) of all elements whose keys hash to the same index. This method is known as **chaining**.




Figure 2.4: Hash table

Pros: Very fast operations (search, insert, delete) with average $O(1)$ time—performance stays the same even if you have 100 or 10 million elements. The fastest data structure for lookups.

Cons: No ordering: data is scattered pseudo-randomly, so there's no way to print elements in order. Efficiency depends entirely on the quality of the hash function—a poor hash produces lots of collisions and slows everything down.